


THE BENCHMARK ON THE STREET AND HOW THE KERB IS SHARED

FIG.10

What a benchmark looks like at eye level

A BENCHMARK AT EYE LEVEL · third-order BM-301 · elevation 1 m = 138 units

a type drawing; it carries no orientation



Two of the five standard components are hard constraints: seating must have armrests, and the complaint entry must offer a non-scan method. A rule you cannot see in the drawing is a rule the drawing does not contain.

The same kerb, after snow clearance

In winter the order becomes one real trade-off

Piled on the device band

Devices stop; S06 falls back to staffed transfer

Piled on the pedestrian band

Pedestrians pushed into the carriageway - the worst of the three


Piled on the kerbside parking band

The one item in the order that may be lost first

This proposal takes this one

Kerb allocation in priority order

KERB ALLOCATION IN PRIORITY ORDER · sequence and precedence only, no widths



The order is itself a position: device charging sits behind pedestrians and accessibility, which means introducing devices must not be paid for with degraded walking conditions. Winter snow storage must be reserved at this stage alongside the device and pedestrian lanes — piled on the device lane the devices stop; piled on the pedestrian lane people are pushed toward the carriageway, which is worse.

The elevation is drawn at 1 m = 138 units. The figures are a scale reference and a statement about who is present, not a rendering. The section expresses sequence and precedence only and carries no widths — capacity must be computed on site from measured clear width, and a figure not derived from measurement is fabricated certainty. Component specifications are in geometry/public_space.geojson and the scenario cards; positions await official boundaries and title verification.

THE LEVELING LINE · Haidian, Beijing · jiangmuran · Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-19 (sources.json; 1 not obtained)

Plan EPSG:4548/4326 · heights relative; the 1985 national datum is not used

LANDMARKS, KIT OF PARTS, SIGNAGE SYNTAX AND OPERATING CYCLE

FIG.09

a type drawing; it carries no orientation

Three AI pilgrimage landmarks

THREE LANDMARKS · spoken in engineering language; no spectacle

L1 The origin benchmark stone


BM-0 · AI Origin Community



Every merged proposal's number is set into the paving — contributors' GitHub IDs are inscribed here. It is not grand. It is accurate, and must be accurate enough to be reused a century later.

L2 The closure stele

BM-1 · Zhongzhiyuan



A public reading wall, updated live, showing each scenario's closure error against its tolerance. Exceedances are marked in datum red. The value of this landmark is that it is allowed to look bad.

L3 return-to-datum point


BM-2 · Dazhongsi



An annual civic ceremony space: the line's readings are zeroed here, tolerance revisions are read out, and scenarios returned for re-survey in the past year have their disposition explained. On ordinary days it is a public dwelling space.

Kit of parts for public space: five standard components

KIT OF PARTS · common specifications and open drawings, so any new node joins in one language



Signage numbering syntax (agent.5)

A visitor reading the number knows which order they are standing at and how often it is re-measured

BM-0	Origin benchmark	Annual origin and closure computation
BM-1 / BM-2	First-order benchmark	Annual re-survey
BM-2x	2nd-order point	Quarterly re-survey
BM-3xx	3rd-order point	Monthly re-survey
RT-N / RT-S	Closing route	Semi-annual re-survey of the whole line

Operations are organised by re-survey cycle, not by festival calendar (agent.6)

Events are therefore governance actions, not publicity

Monthly	Community day	third order BM-301/BM-302/BM-303	P4 · P5 · P7
Quarterly	Scenario open day	second order BM-21/BM-22	P2 · P3
Semi-annual	Route re-survey	first order, routes RT-N / RT-S	P1 · P2
Annual	Zeroing ceremony	L3 return-to-datum point	all four review parties present

Developer conversion path: take part in re-survey → propose a scenario → enter the test field → obtain closure clearance → operate
International communication uses published readings as its material, never commitments.

The f=0.26 / f=0.14 on the reading boards are illustrative, not measured results. f is always a deviation: larger is worse, and f ≤ F is the passing test. Landmark siting requires heritage and engineering review; this proposal does not pre-empt it. The components are a directional proposal for a professional team to develop and are not construction drawings.

THE LEVELING LINE · Haidian, Beijing · jiangmuran · Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-19 (sources.json; 1 not obtained)

Plan EPSG:4548/4326 · heights relative; the 1985 national datum is not used

REGIONAL COORDINATION INTERFACES: EXTENDING THE LEVELLING NETWORK

FIG.11

An executable form of coordination: how two levelling networks join



A joint observation produces a number, not an intention:

Each network reads the same public question set; the difference is the cross-network closure error. Within a jointly agreed tolerance, the two networks' records may be mutually recognised; over it, each re-surveys, and no declaration of partnership substitutes for a reading.

Nobody gives up authority: a joint observation asks only that conventions be comparable, not that organisations merge.

a type drawing; it carries no orientation

What each of the five partners exchanges, and what it does not

Partner	Stage	Exchanges	Explicitly does not exchange	Precondition for the interface
BDA (Yizhuang)	Volume production and high-level AV road testing	Mutual recognition of closure records across complementary speed domains	Not one admission: an F1 pass here admits a device to pedestrian density, not to a carriageway, and the reverse also holds	Both parties confirm the speed-domain split and test protocols
Future Science City	Corporate research institutes and engineering	Scenario demand and real user density	No claim on IP; no scheduling on anyone's behalf	Wing benchmarks built and a first baseline taken
Huairou Science City	Large facilities and basic research	The measurement method itself: how f is defined, cycle length, the statistics behind revising F	No scenario data is exchanged; the coordination is methodological, not applied	The closure convention published and open to methodological review
Beiwei Community	The innovation community the taskbook names	Reciprocal joint points; exchanged review parties	This proposal does not know its positioning and does not set its benchmark order	Each party names at least one publicly reachable joint point
Beijing-Tianjin-Hebei	Cross-provincial coordination	Cross-provincial recognition of the tolerance convention	No direct recognition of records; without a shared convention they are not comparable	Tolerance definition and revision procedure agreed across provinces

None of the five rows has been confirmed by the party named.
No internal plans are known, nobody is committed, no agreement is assumed; each row is an interface to evaluate independently.

Each partner's stage is described from publicly available general positioning and is unconfirmed by them. The speed-domain relationship in the first row must be settled against both parties' actual test protocols. This sheet proposes interfaces; it commits nobody and is no government determination.

THE LEVELING LINE • Haidian, Beijing • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-19 (sources.json; 1 not obtained)


Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used


THE THREE KEY AREAS, THE TWO WINGS AND THE BENCHMARK LAYOUT

FIG.04

Each area and wing carries a survey role — origin, tolerance, density; the blue dashes are the ten-minute walk catchment, by network distance, not radius.

orientation is noted on each drawing






EVIDENCE CHAIN AS AN UNCLOSED LEVELING CIRCUIT

FIG.02

a type drawing; it carries no orientation



Act One: measuring this open call itself

MEASURING THE OPEN CALL ITSELF · as of 2026-08-13 · 793 merged · 793/793 fetched · PR numbers past 1000

793

Merged proposals (git tree, authoritative)

The gallery index lists 507 at the same moment — 287 behind

18.2%

Machine-readable disclosure field empty

144 of 793 still hold the scaffold placeholder or are empty

228 → 9

Distinct strings written → actual model families

The GPT/Codex family alone is written 85 ways across 355 proposals

The “machine-readable” disclosure field is not machine-readable.

charter.6 requires disclosing the generation method and charter.5 requires it structured and agent-readable, yet none of the four gates checks this field. Nobody can aggregate “which models produced this call” from it.

The fix is light: split “model” into an enumerated family plus free-text detail, and add one enumeration check.

This proposal is the instrument for the missing last step: compute the closure error, and give it consequences.

Registers answer “what did the system declare” and assessments answer “what do experts think”; closure error answers “how much does the conclusion differ across nodes and across time for the same public question”.

Sources: visual/assets/census.json, field_map.json and the first-round (184-file) snapshot agent_declarations.json, all shipped with the package. The authoritative source is the GitHub git tree, not the gallery index, which lags. Generation scripts are in the accompanying issue and are re-runnable. The corpus grows daily; recompute before citing.

THE LEVELING LINE · Haidian, Beijing · jiangmuran · Concept advice on provisional boundaries; not a substitute for statutory planning

Data to 2026-08-19 (sources.json; 1 not obtained)

Plan EPSG:4548/4326 · heights relative; the 1985 national datum is not used

RECOMPUTED METRICS AND CLOSURE EVIDENCE

FIG.06

a type drawing; it carries no orientation

Run once: 10 cases, 2 accepted, 8 refused, reasons quoted from shipped files; 12 cards name the path that remains with the AI off.

Concept advice: the boundaries are provisional substitutes, and every digit is computed from the submitted geometry — not surveyed precision.

METRICS RECOMPUTATION

RECOMPUTED METRICS · EPSG:4548

Class 1, measured: recomputable directly from this package's geometry

Overall design area	11,412,825 m ²	S06 low-speed device pilot	237,479 m ²
Spine length	9,443 m	S10 public-safety review	15,369 m ²
Benchmark count	8	Green ratio (land-use partition basis)	0.1220
Green corridor ratio	0.1966	Phase increments, summed	2,920,311 m ²
Public measurement-point ratio	0.0642	Phased share of design scope	0.2559
Indicative footprint area	82,276 m ²	Existing street length, measured	157,510 m
Green corridor area (footprints deducted)	2,244,082 m ²	Existing network density, m/m ²	0.0138
Public measurement-point area	733,011 m ²	Existing median block, measured	4,300 m ²
Key areas	3	Existing building coverage, lower bound	0.1798
Key areas, total	3,692,893 m ²	Six spine benchmarks walked, mapped net	11,472 m
S11 industry test ground	53,008 m ²		

A further 32 count metrics (unit: count) ship in metrics.json: cards, personas, landmarks, cases, projects, phases, land-use classes, errata — all recomputed by the build, none in this table.

Class 2: requires official regulatory support; held at unknown

Floor area ratio · building height · density · setbacks · road redlines

Filling estimates into a gap while official data is absent is fabricated certainty.

Class 3: requires continuous re-survey calibration; no baseline yet

Per-scenario closure error f · tolerance compliance rate · non-AI path coverage · resident-initiated re-surveys

Baselines follow one near-term cycle. This proposal says no data exists rather than passing intent off as measurement.

Evidence ceiling: computed exactly ≠ measured reliably

Each metric's ceiling is the weakest of the files it reads, computed by the build rather than judged by the author




60 metrics, 59 of them gradeable; 23 are drawn on an inferred boundary and 0 reach the regulatory baseline.

Every class-1 value is recomputed from the shipped layers by the accompanying script;

A number nobody can recompute is not evidence — and that standard applies to this proposal first.

Walk to the nearest benchmark: 9 segments, 6 of them failing

WALK TO THE NEAREST BENCHMARK · worst per segment; red line 540 m (P5 at 0.6 m/s, 15 min)



Worst 1,111 m; for P5 that is 30.9 minutes - this proposal's own rule, and the layout as drawn does not meet it. Closing it needs 9 more third-order points, whose annual operating cost is already in this package's cost table rather than left to a next version.

Sources: this package's metrics.json and visual/assets/field_map.json, both shipped with it; the authoritative source is the GitHub git tree. The independent verifier visual/assets/verify.js checks every class-1 metric. Generation scripts are in the accompanying issue. The corpus grows daily; re-run before citing.

THE LEVELING LINE · Haidian, Beijing · jiangmuran · Concept advice on provisional boundaries; not a substitute for statutory planning

Closure evidence and measured corpus

CLOSURE EVIDENCE AND MEASURED CORPUS · n=793 (git tree) · 2026-08-13

Closure evidence: the mechanism has been run

`node visual/assets/run\_s08\_tabletop.js` · 10 cases · 2 accepted · 8 refused · re-runnable by the reviewer · rule and reason quoted verbatim from the shipped file

- accept

f <= F is level for this cycle and the scenario continues
- accept

f > F returns the whole route, not the failing station
- refuse

rule 5: amending only the worst station is not an available action
- refuse

record.stations: needs at least 3 items, has 2; review parties: 2 distinct across the circuit; a reading taken by one kind of party cannot detect what that kind of party cannot see
- refuse

a record does not get to state its own f
- refuse

verdict.f declares 0.13 but the stations recompute to 0.360000; verdict.passed says true but f 0.360 against F 0.2 derives false; verdict.action is “continue” but f against F requires “return-whole-route”
- refuse

a circuit read by one kind of party cannot detect what that party cannot see
- refuse

review parties: 1 distinct across the circuit; a reading taken by one kind of party cannot detect what that kind of party cannot see
- refuse

without an equivalent non-AI path the stop rule would be circumvented
- refuse

record.verdict.non_ai_path_available: must be true, got false; verdict.non_ai_path_available is false: a scenario with no non-AI equivalent cannot be returned, because returning it would interrupt public service — so the stop rule would be circumvented
- refuse

omitting the non-AI path field is not the same as satisfying it
- refuse

record.verdict: missing required property “non_ai_path_available”; verdict.non_ai_path_available is absent: the equivalent non-AI path is not waivable, so omitting the field cannot be treated as satisfying it. State it explicitly as true.
- refuse

a connecting route must terminate at a first-order benchmark, or nothing checks the carried value
- refuse

stations(last): a connecting route must terminate at a first-order benchmark, because the carried value has nothing to be checked against otherwise. This one ends at BM-302 (third).
- refuse

rule 7: two consecutive passes are required, because once may be luck
- refuse

resumption needs two consecutive passing cycles, record shows 1; once may be luck
- refuse

rule 7: an F1 incident resumes only on unanimous confirmation, no majority rule
- refuse

resumption after an F1 incident requires unanimous confirmation by all four review categories

Nothing here is a field measurement. Not one reading was taken: these are worked values chosen to exercise the refusal branches.

Measured corpus: few take this road


793 packages, read from the git tree (507 listed in the index). Pulse or lifeline is at 168 and proof or verification at 184, both saturated; measurement or levelling is at 10. Motif detection is a keyword regex and therefore a lower bound, and the convergence on three-areas-two-wings is induced by the taskbook rather than agreed.

Leveling is not another railway-operations metaphor. It is a method for how trust is established.

What is left of the belt with the AI switched off


Without a non-AI equivalent the stop rule is circumvented. All 12 cards name their equivalent; the 4 F1 ones keep theirs permanently (full text in scenario_cards.json).

S01	F3	Paper sign-in and on-site stewards	S02	F2	Residents report breaks on paper at the nearest third-order benchmark...	S03	F2	The staffed counter runs independently of the agent
S04	F1	A staffed help desk and a printed directory of institutions, departments and clinic hours	S05	F2	Authorisation credentials and change records are documents	S06	F1	Devices stop and return to their bays
S07	F3	The evidence hall's physical display and paper archive stay open	S08	F3	Physical interpretation plates along the route and scheduled human-led tours	S09	F2	Prices and licences go on a physical notice board
S10	F1	A review meeting and a paper record	S11	F1	The test ground stops serving the public and its boundary is held by physical...	S12	F2	Checking reverts to user reports on paper or by phone plus quarterly...



THE LEVELING LINE

A city that publishes its own error.



Leveling is not about measuring accurately. It is about measuring back.

Depart from the origin, carry the height station by station, run the circuit and return — the part that does not land back on the datum is the closure error.


Within tolerance every station is accepted; over tolerance the whole run is void and re-measured, and no single station may be patched.

This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and AI public services.

OVERALL CONCEPT AND SITE OVERVIEW

Research area agrees 100% with the surveyed park; the gap is in the design area only, nearest 412.5 m — and this proposal's spine is displaced 1,116.7 m with it.

Concept advice: the boundaries are provisional substitutes, and every digit is computed from the submitted geometry — not surveyed precision.



Legend and how to read it

LEGEND • surveyed data and inferred boundaries are shown separately

Surveyed (independently checkable)

- Jing-Zhang Railway Heritage Park, phase-one built segment (OSM survey, 17.49 ha)
- Disused railway (14 segments, 3,028 m)
- Existing street network (OpenStreetMap, 1,376 ways in this window, expressway to tertiary; local lanes not drawn)
- Existing building footprints (OpenStreetMap, 7,709 in this window)

Inferred (not usable as a redline)

- Coordinated research area 43.6 km²
- Overall design area 11.4 km²

This proposal's design (generated against inferred boundaries)

Blank-block edges follow the existing arterials; they are not a parcel subdivision or an ownership boundary

- Cultural use (0803)
- AI R&D and research (0802)
- Community services and support (0702)
- Industry and commercial services (09)
- Park / green (1401)
- Urban road land (1207)
- Left blank by this proposal (code 16)
- Spine (design centreline)
- Tiered benchmarks

The line to read on this sheet is the red one.

Research area and surveyed park agree 100%, so the announced bounds do not conflict with actual geography; the discrepancy is confined to the overall design area. This proposal's spine is displaced with it. Marked here rather than hidden.

What the belt does

WHAT THE BELT DOES • the five functions are five positions on one loop

- Set the tolerance
- Depart
- Take the readings
- Close at the origin
- Over tolerance: re-survey

BM-1 Zhongzhiyuan
BM-0 Origin Community
BM-2 Dazhongsi, the wings
BM-0
the whole route returns

Over tolerance the whole route returns for re-survey and drops to the non-AI equivalent - the loop does not close, and the belt does not run.

9.4 km
Spine length

11.4 km²
Overall design area

8
Tiered benchmarks (four orders)

Site cross-check (the red line at left):

The announcement draws the design area 1-2 km around the Jing-Zhang Rail Heritage Park - public ground already open. That park, from OSM, is 17.49 ha: 100% in the research area, 0% in the design area, nearest 412.5 m. Same instrument on this proposal: the spine sits 1,116.7 m from that park - the distance to be brought home when the official polygons publish. A re-computability standard invoked only when it flatters the author is not a standard.