

Haidian District, Beijing • both sides of the Jing-Zhang Rail Heritage Park • three key areas: Zhongzhiyuan • Beijing AI Origin Community • Dazhongsi AI Industry Cluster

A civic AI leveling network that uses closure error as its trust metric

Leveling is not about measuring accurately; it is about measuring back.

Depart, carry the height station by station, return — that gap is f.

Within tolerance, every station on the line is accepted;

over tolerance the whole run is void; no station may be patched alone.

That hundred-year-old rule answers this belt's hardest question:

what a city's trust in its AI rests on: not performance, but return.



Ground within a ten-minute walk of a benchmark

Dashed: 232 buildings today. Filled: the ground that only comes within ten minutes once the eleven stitching links and three spurs are built — 261 in all, a gain of 29.

Catchment 203.3 ha to 220.6 ha; 800 m of network distance, not a radius.

READINGS, ALL RECOMPUTED FROM THE SHIPPED FILES

11,412,825 m²

Overall design area

9,443 m

Spine length

8

Tiered benchmarks

13.8 km/km²

Network density, measured

17

Oversized blocks, not parks

19.7%

Green corridor, share of design area

What this proposal does not give here matters as much as what it does

Plot ratio, height, coverage, setbacks, road redlines — inside the statutory approval process; none is drafted here before the official conditions are published.

The tolerance values per scenario — set by the tolerance assembly at BM-1; this proposal gives the grading basis and drafts no value.

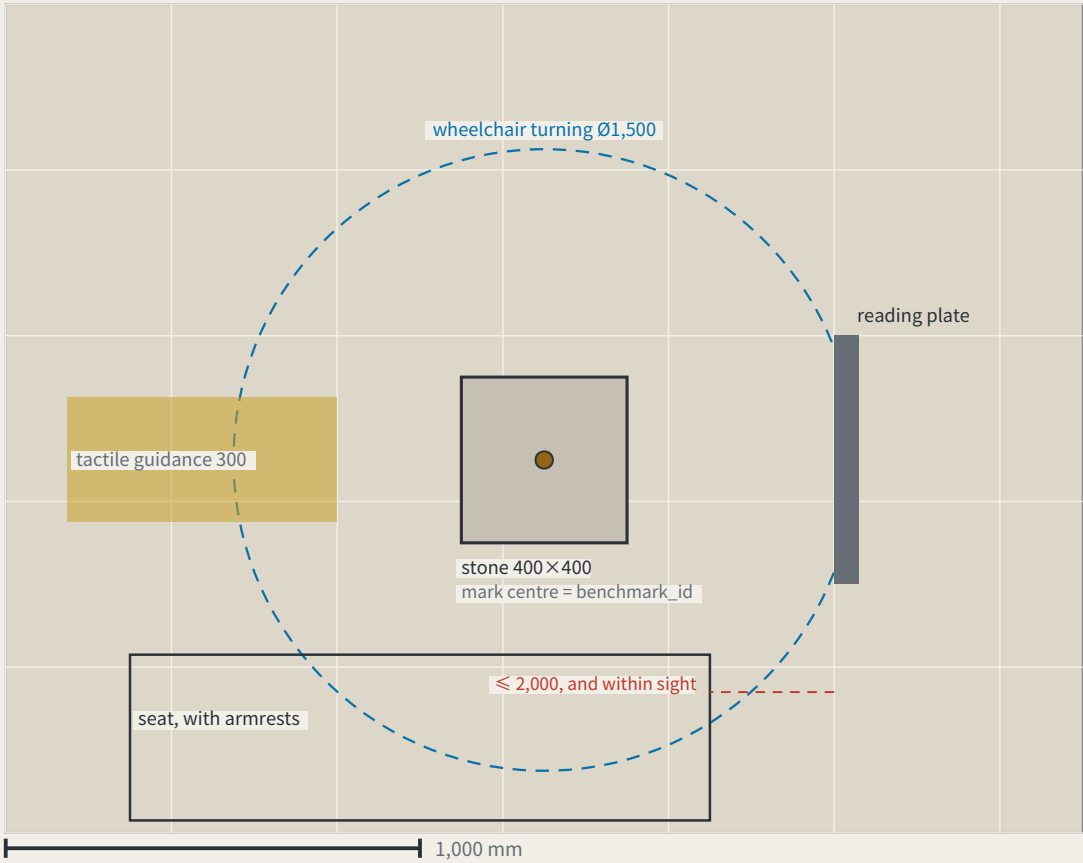
Any engineering feasibility conclusion — stitching points register a need for connection; transport and structural review is outside this proposal.

FIG.00	The Leveling Line: a city that publishes its own error		FIG.17	The step-free link from the origin community to the heritage park	10
FIG.01	Overall concept and site overview	6	FIG.18	Maintenance: putting the wear on the parts that can be swapped	4
FIG.02	Evidence chain and submission package: a leveling circuit not yet cl...	6	FIG.19	Monuments beside heritage fabric: the no-drilling construction and ...	10
FIG.03	Land-use structure and the three scope levels	6	FIG.20	Reading after dark, without lighting the benchmark	10
FIG.04	The three key areas, the two wings and the benchmark layout	7	FIG.21	How far the nearest benchmark actually is: this proposal's own rule...	10
FIG.05	Slow mobility, blue-green and closing routes	7	FIG.22	Winter: water, ice, and two of this package's rules in collision	11
FIG.06	Recomputed metrics and closure evidence	7	FIG.23	Finding the nearest benchmark: which way, and how far	11
FIG.07	Identity: mark, construction and applications	7	FIG.24	The device envelope: what exactly goes in those 18 m ²	11
FIG.08	Innovation ecosystem and element mechanisms	8	FIG.25	BM-0, the origin: the place both routes have to return to	11
FIG.09	Landmarks, kit of parts, signage syntax and operating cycle	8	FIG.26	BM-2x, the second order: a scenario has to be stoppable, and that t...	12
FIG.10	The benchmark on the street, and how the kerb is shared	8	FIG.27	One reading: where a person stands, and whether they block the way	12
FIG.11	Regional coordination interfaces: extending the levelling network	8	FIG.28	Third order: the most numerous tier, and more than one reader	12
FIG.12	Four-quadrant pedestrian connection and device queue reservoir a...	9	FIG.29	The annual first-order closure: how long a route actually takes	12
FIG.13	The three key areas in section, at one scale	9	FIG.30	The year: where forty-seven readings fall	13
FIG.14	The through-block public route to the Qing river	9	FIG.31	The other year: the hours this network asks of people it does not p...	13
FIG.15	Phasing: advanced by closure results, not by dates	9	FIG.32	How many people: the cheapest roster is the one that empties the inst...	13
FIG.16	The benchmark and its reading plate, as a construction detail	3	FIG.33	Dazhongsi: where the belt meets ground this network does not own	13

FIG.16

a type drawing; it carries no orientation

I. Plan

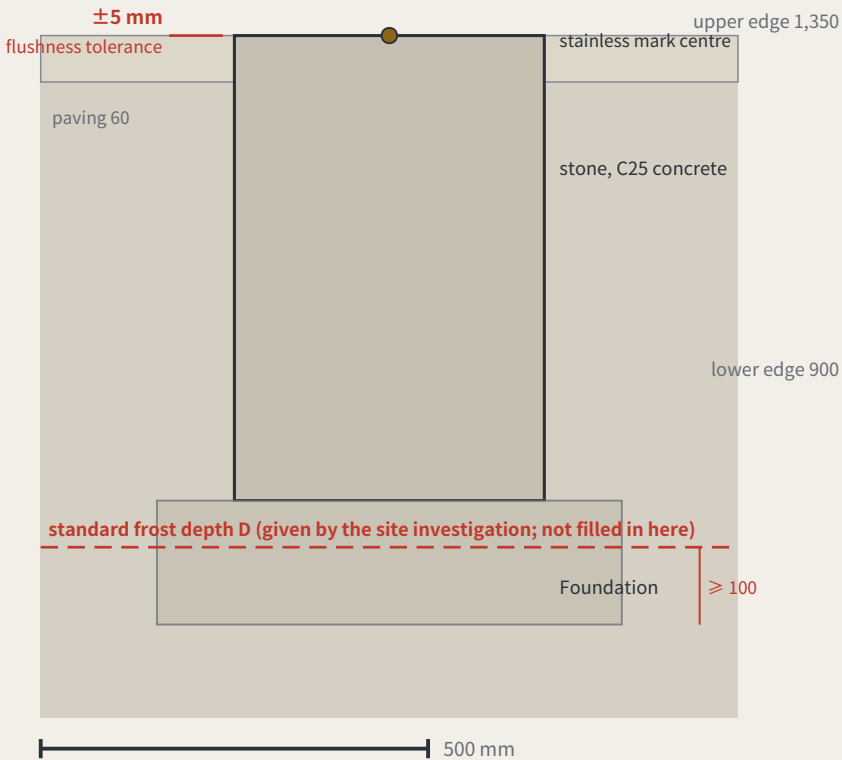


Dimensions: design intent and as built

The right column is deliberately empty — as-built values are measured on site


Item	Design intent	Why this number	As measured
The stone	400 × 400 × 600 mm, C25 concrete, stainless mark centre in the top face	the mark centre is the point a benchmark_id refers to	
Stone top against the paving	≤ ±5 mm	KIT-01 says flush with the ground and no trip hazard; without a tolerance, “flush” is an opinion	
Depth to underside of foundation	≥ the local standard frost depth D + 100 mm	D comes from the site investigation and is not filled in here — this package holds no verifiable frost-depth source for Haidian	
Plate face size	600 × 450 mm, raked 15°	the rake cuts glare and standing water; the face is replaceable without the post	
Plate lower / upper edge	900 mm / 1,350 mm	the band a seated and a standing reader share; P5’ s continuous step-free route should not end at a plate they cannot read	
Appeal panel	QR code, phone number and in-person address, all three side by side	KIT-05: a QR code alone excludes anyone without a smartphone from appealing, which makes persona P4 unworkable	
Tactile guidance strip	300 mm wide, stopping 300 mm short of the stone	the stop line leaves room to stand and take a reading rather than leading onto the stone	
Wheelchair turning space	Ø1,500 mm, headroom ≥ 2,100 mm	the turning space coincides with the reading position, so taking a reading needs no reversing first	
Seat	seat 450 mm high, with armrests, within 2,000 mm of the plate	KIT-03: taking a reading should not require standing and waiting, and the seat must be in sight of the plate	

II. Section



III. Plate elevation (post shown broken)

Specimen values; F is a dimensionless rate, read from scenario_cards.json




the face panel is replaceable without touching the post

The whole proposal rests on one object, and not one dimension had appeared anywhere in the package — a component library without dimensions is a mood board. Here it is in plan, section and elevation, nine dimensions each with the reason it is that number; read the scale off the bars. The one figure left blank is the frost depth D: no verifiable source here, so the sheet gives the rule and not the number. The as-built column is empty too. Concept design intent; requires detailed design and engineering acceptance.

a type drawing; it carries no orientation


I. Order of removal: from what is swapped to what is never moved



One person is assumed to carry up to 23 kg. The stone, 400 × 400 × 600 mm of plain concrete at 2,400 kg/m³, works out at about 230 kg — the one part nobody can carry is exactly the one part that should never be replaced.

IV. What these intervals add up to

Summing the intervals above across 8 benchmarks: about 1.42 replacements per point per year, roughly 11 across the line, about 6 staff-hours a year. Those hours were not in the annual operating table — it counted re-survey sessions only. They are now: paid hours go from 133–214 h to 139–220 h, adding about CNY 680–1,473 a year. Maintenance is not a hidden cost, but it genuinely was not in the table.



Intervals and hours per swap are assumptions, not measurements; the stone is never replaced and is excluded.

II. What fails, how often it is replaced, and whether it touches the datum

Part	How it fails	Replaced	Mass	Tool	Touches datum
Face panel of the reading plate	stale values, scratching, fading	checked each cycle, replaced when damaged	3.0 kg	hex key	no
QR and appeal sticker	the link dies, the surface wears	annual	0.2 kg	none	no
post	impact, corrosion	about 20 years, or after an impact	12.0 kg	socket wrench	no
seat panel	wear, graffiti	about 10 years	18.0 kg	socket wrench	no
Tactile guidance strip	abrasion, and repaving	about 15 years, with the paving	6.0 kg	paving tools	no
The stone	impact, excavation, settlement	never replaced	230 kg	plant required	yes

III. After it fails: found → held → replaced → lost

Found	any re-survey that finds the face stale or missing records it in that cycle the plate’ s own rule — a stale value counts as not posted — is the detector
Held	the point is marked returned in datum red, shown at every third-order point in the area bad news belongs where it happened (errata E50)
Replaced	a standard part is swapped without touching the stone or its mark centre one person, one visit; heaviest part 18 kg
Stone lost	the height is re-established from the next order up; the closure record carries the break; the id is retired and never reused reusing an id would quietly join two different points into one history

Method: trust comes from closure, not from one accurate reading

Leveling is not about measuring accurately; it is about measuring back. Depart from a known point, run the circuit, return — the difference is the closure error. Over tolerance the whole run is void and re-measured, and you may not patch the worst station and keep the rest.

This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and autonomous shuttles, and AI health, education, legal and daily services. Urban AI governance is the method layer here, not the selling point.

Two prohibitions at the heart of the mechanism

One: amending only the worst station is not allowed. This makes tuning until the metric looks good structurally ineffective.

Two: tolerance may tighten on evidence and may never loosen because a scenario failed. This closes off the moving goalpost.

f is always defined as a deviation, never an attainment score: classification scenarios take 1 minus the consistency ratio, service scenarios the satisfaction range, safety scenarios the sum of false-positive and false-negative rates. All three run the same direction, so $f \leq F$ is always the passing test.

Act One: measure this open call first

228	Merged proposals (git tree, authoritative)
30.3%	Machine-readable disclosure field empty
84 → 8	Distinct strings written → model families

The “machine-readable” disclosure field cannot actually be aggregated: the GPT/Codex family alone is written 44 ways across 103 proposals. The fix is light — split the field into an enumerated family plus free-text detail,

and add one check to the gates.

Closure error in the site's own data

412.5 m	Provisional design area to surveyed park
100%	Agreement with the research area

Checking the inferred boundary against OpenStreetMap's surveyed heritage park, two independent routes to the same object do not close. The 43.6 km² level agrees completely; the discrepancy is confined to the 11.4 km² level.

The same measurement also measured this proposal: the submitted spine lies 1,116.7 m from the surveyed park, because it was generated against the provisional boundary. This is stated because a recomputability standard invoked only when it flatters the author is not a standard.

Main front one: low-speed robots and autonomous shuttles

No benchmark, no robot. This turns governance into a spatial design problem: to expand the operating area you must first build measurement points.

Across the field, speed limits, remote and physical emergency stop, on-site safety officers, incident logs, an equivalent non-AI path and scenario-level exit are already the de facto standard. This proposal adopts all of them but presents none as an innovation — they are the admission floor.

The increment is where nobody has been: ice and low-temperature re-survey (zero hits across the field), noise as a number (zero hits), jurisdictional seams, fleet density ceilings, and emergency-access yielding. Their shared structure is that the same system behaves inconsistently across conditions or jurisdictions — which is exactly what closure error measures.

Any safety incident suspends that entire machine type network-wide, not the individual unit. Tolerance scales with kinetic energy: faster or heavier requires a stricter closure clearance first, never an exemption.

Jurisdictional seams: where pilots actually die

8 / 8	Cross-jurisdiction points on this belt
-------	--

Crossing jurisdictions is the normal condition, not an edge case. The failure mode at a seam is not a contest over authority but that each side reasonably concludes it is not theirs, so the device keeps running unreviewed until something happens. The rule: if neither side holds a valid reading, the section is closed — so the default consequence of inaction is that the device stops.

Main front two: AI public services

Do not measure the average; measure the dispersion. The risk is not in mean accuracy but in how much the conclusion differs when different people ask the same question — which is precisely what closure error measures.

Three non-waivable boundaries: prescriptive judgements must be made by qualified people and logged; the equivalent non-AI path is permanent and must not be slower; and no profiles of identifiable individuals and no cross-scenario linkage.

Statement of boundaries

Everything here is open collaborative concept advice on provisional substitute boundaries. It does not replace statutory planning and constitutes no government determination, approval basis or implementation commitment. No floor area ratio, building height, density, setback or road redline figures are given, and no heritage protection zone or blue line is inferred. Responsibility is written as role types only, all subject to negotiation. Final judgement rests with people and professional teams.

A city that publishes its own error.



FIG.01

Concept advice: the boundaries are provisional substitutes, and every digit is computed from the submitted geometry — not surveyed precision.

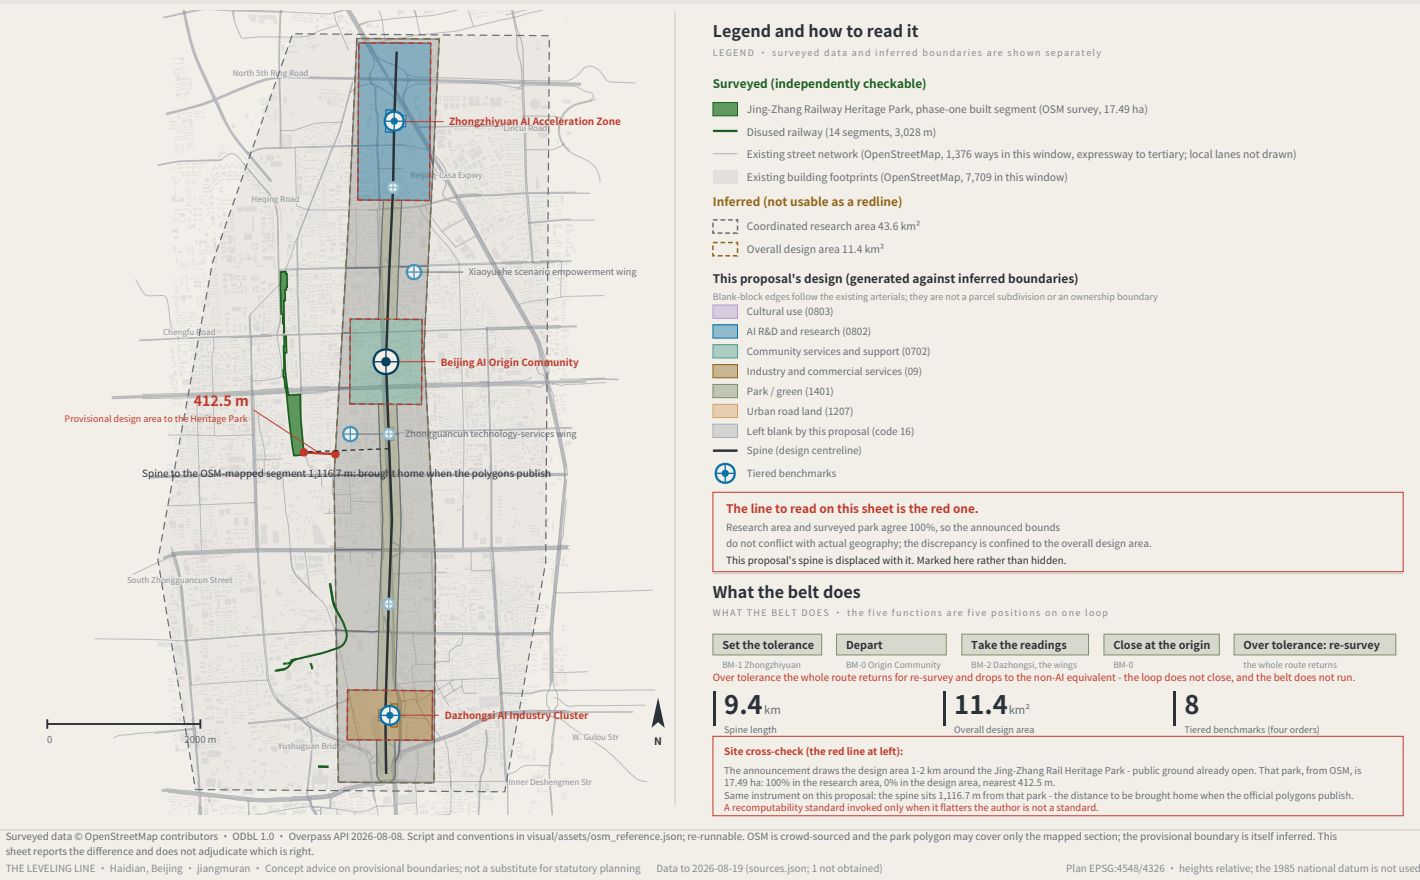



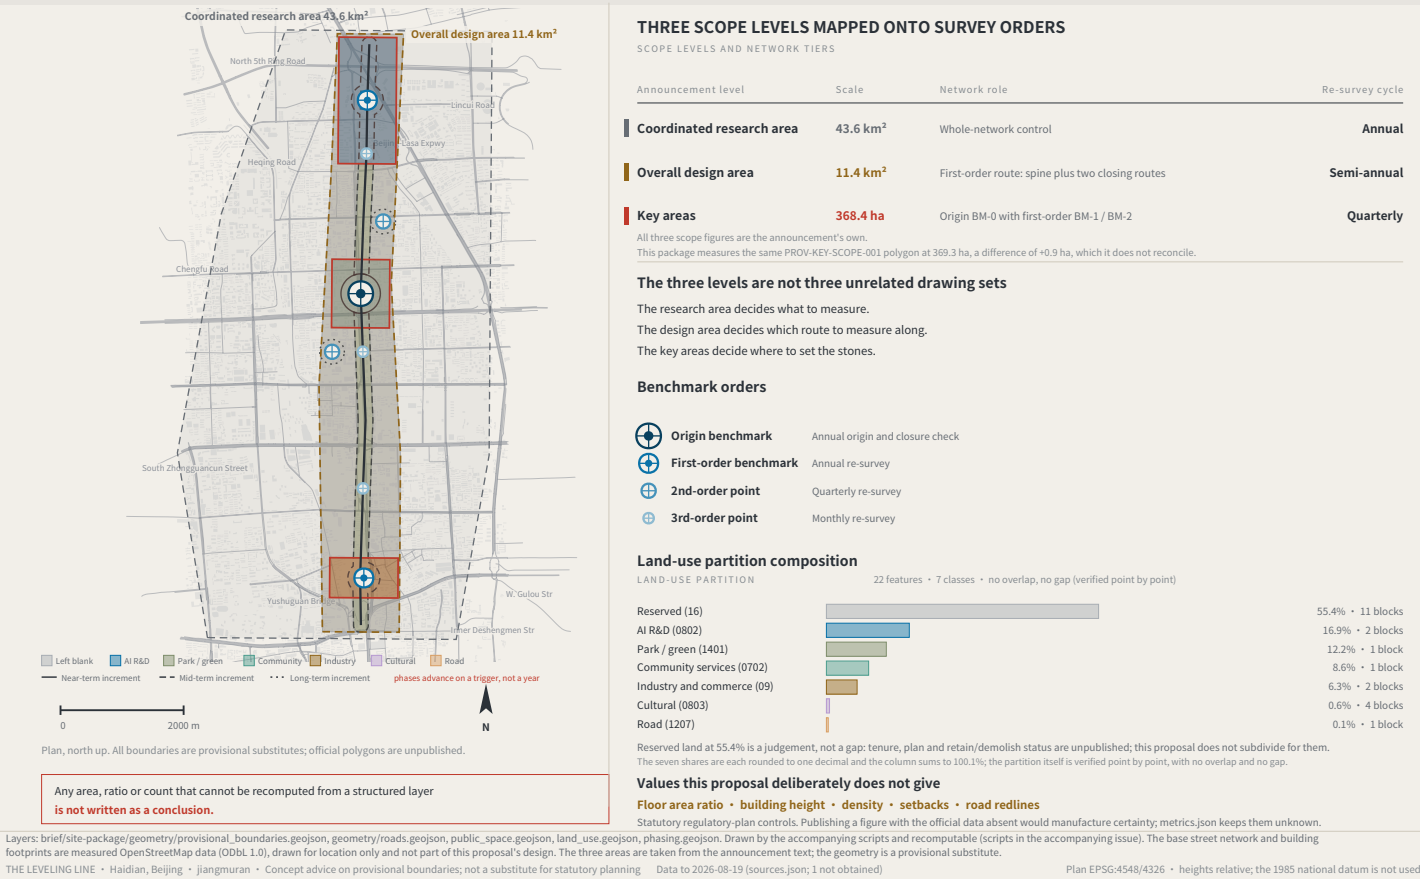
FIG.02



Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used

FIG. 03

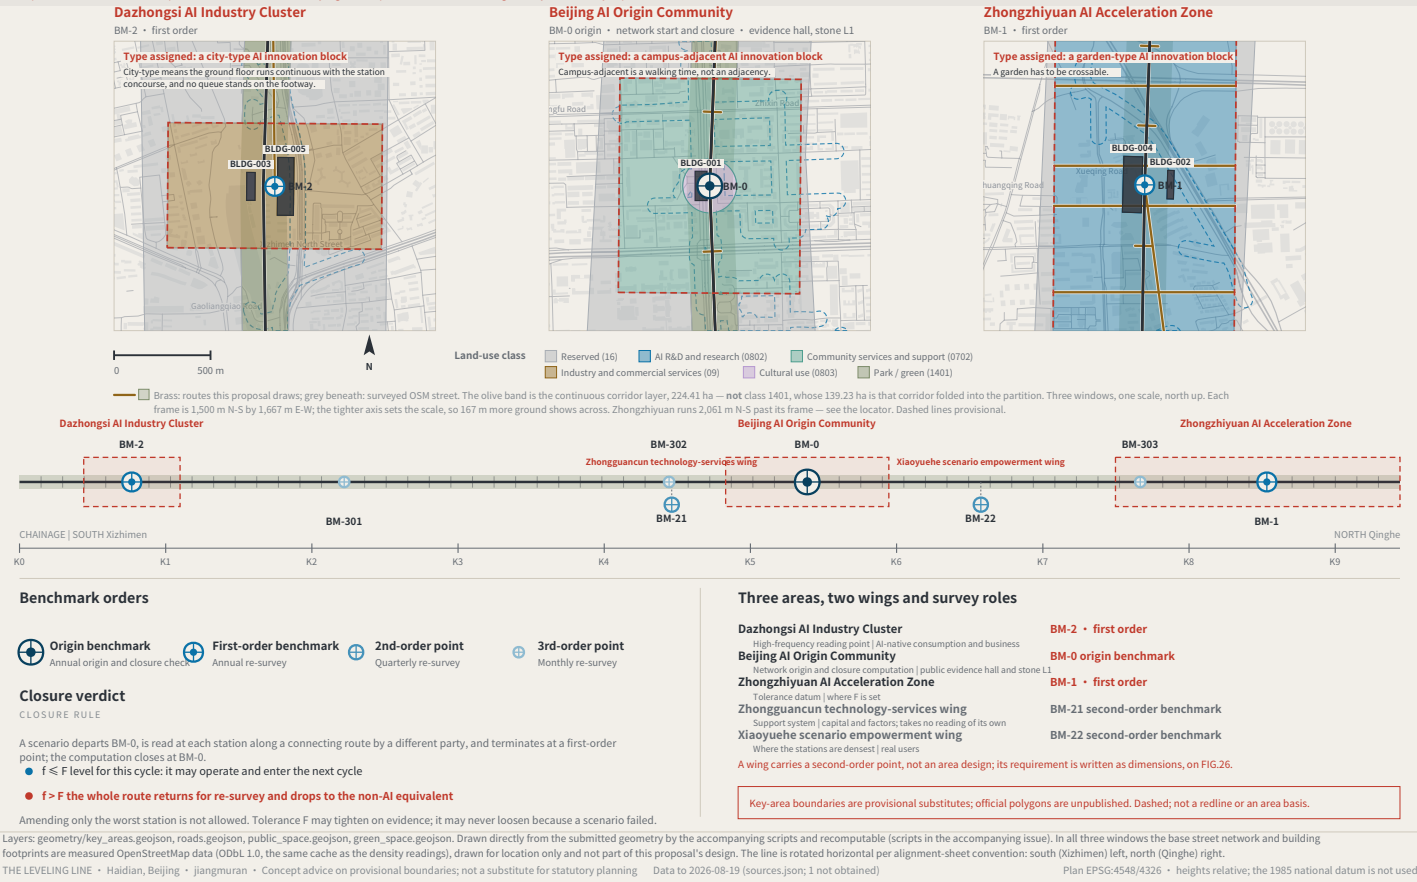
Concept advice: the boundaries are provisional substitutes, and every digit is computed from the submitted geometry — not surveyed precision.



THE THREE KEY AREAS, THE TWO WINGS AND THE BENCHMARK LAYOUT

FIG.04

Concept advice: the boundaries are provisional substitutes, and every digit is computed from the submitted geometry — not surveyed precision.



RECOMPUTED METRICS AND CLOSURE EVIDENCE

FIG.06

Run once: 10 cases, 2 accepted, 8 refused, reasons quoted from shipped files; 12 cards name the path that remains with the AI off.

a type drawing; it carries no orientation

METRICS RECOMPUTATION

RECOMPUTED METRICS • EPSG:4548

Class 1, measured: recomputable directly from this package's geometry

Overall design area	11,412,825 m ²	S06 low-speed device pilot	237,479 m ²
Spine length	9,443 m	S10 public-safety review	15,369 m ²
Benchmark count	8	Green ratio (land-use partition basis)	0.1220
Green corridor ratio	0.1966	Phase increments, summed	2,920,311 m ²
Public measurement-point ratio	0.0642	Phased share of design scope	0.2559
Indicative footprint area	82,276 m ²	Existing street length, measured	157,510 m
Green corridor area (footprints deducted)	2,244,082 m ²	Existing network density, m/m ²	0.0138
Public measurement-point area	733,011 m ²	Existing median block, measured	4,300 m ²
Key areas	3	Existing building coverage, lower bound	0.1798
Key areas, total	3,692,893 m ²	Six spine benchmarks walked, mapped net	11,472 m
S11 industry test ground	53,008 m ²		

A further 32 count metrics (unit: count) ship in metrics.json: cards, personas, landmarks, cases, projects, phases, land-use classes, errata — all computed by the build, none in this table.

Class 2: requires official regulatory support; held at unknown

Floor area ratio • building height • density • setbacks • road redlines

Filling estimates into a gap while official data is absent is fabricated certainty.

Class 3: requires continuous re-survey calibration; no baseline yet

Per-scenario closure error f • tolerance compliance rate • non-AI path coverage • resident-initiated re-surveys

Baselines follow one near-term cycle. This proposal says no data exists rather than passing intent off as measurement.

Evidence ceiling: computed exactly $\neq$ measured reliably

Each metric's ceiling is the weakest of the files it reads, computed by the build rather than judged by the author




60 metrics, 59 of them gradeable; 23 are drawn on an inferred boundary and 0 reach the regulatory baseline.

Every class-1 value is recomputed from the shipped layers by the accompanying script;

A number nobody can recompute is not evidence — and that standard applies to this proposal first.

Walk to the nearest benchmark: 9 segments, 6 of them failing

WALK TO THE NEAREST BENCHMARK • worst per segment; red line 540 m (PS at 0.6 m/s, 15 min)



Worst 1,111 m; for PS that is 30.9 minutes - this proposal's own rule, and the layout as drawn does not meet it. Closing it needs 9 more third-order points, whose annual operating cost is already in this package's cost table rather than left to a next version.

Sources: this package's metrics.json and visual/assets/field_map.json, both shipped with it; the authoritative source is the GitHub git tree. The independent verifier visual/assets/verify.js checks every class-1 metric. Generation scripts are in the accompanying issue. The corpus grows daily; re-run before citing.

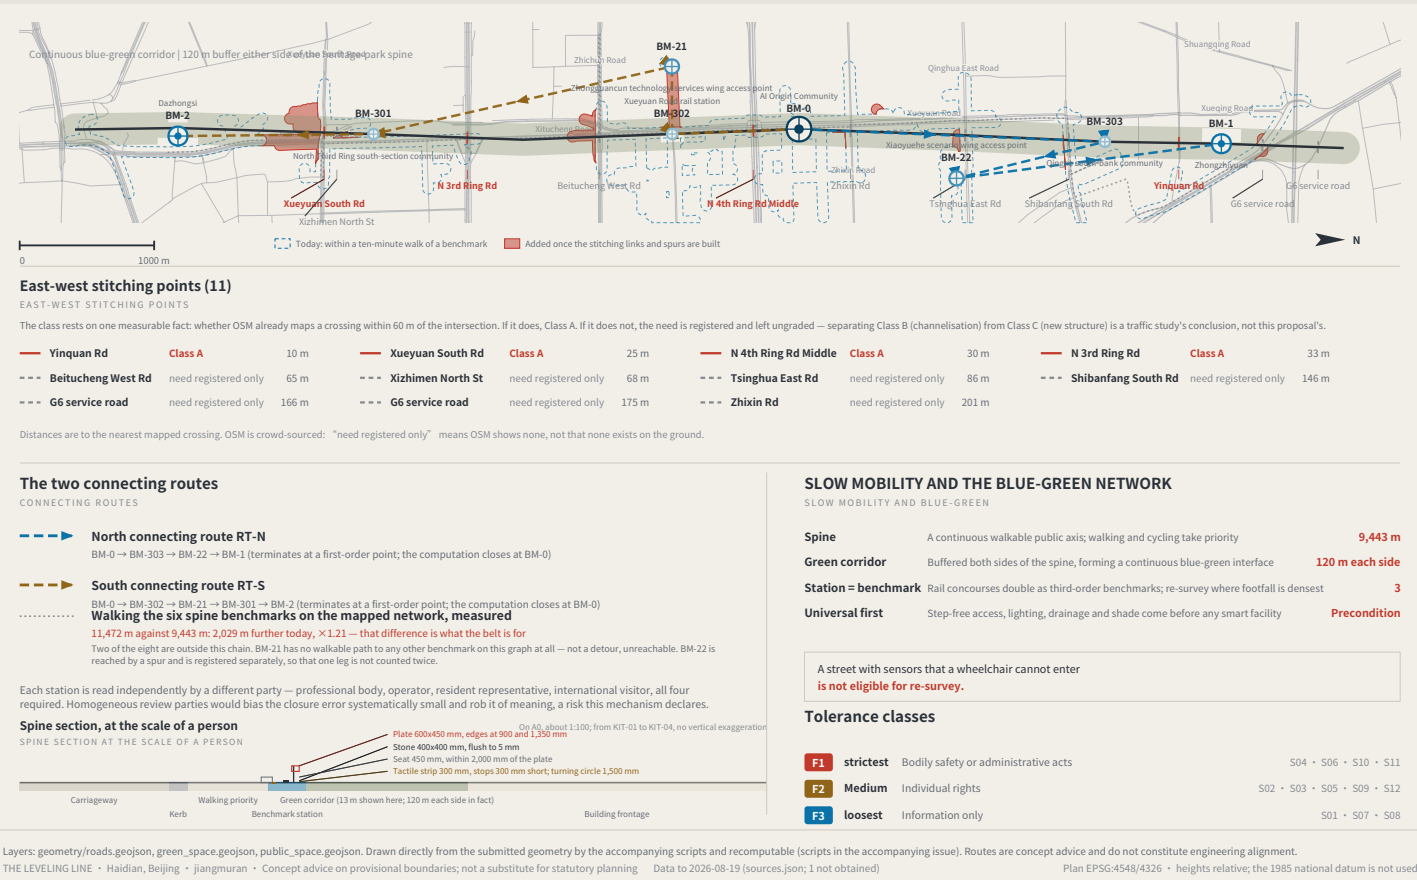
THE LEVELING LINE • Haidian, Beijing • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-19 (sources.json; 1 not obtained) Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used

SLOW MOBILITY, BLUE-GREEN AND CONNECTING ROUTES

FIG.05

Six spine benchmarks: 11,472 m walked today against 9,443 m drawn (x1.21); with the stitching and spurs built, buildings within a ten-minute walk rise 232 to 261.

Concept advice: the boundaries are provisional substitutes, and every digit is computed from the submitted geometry — not surveyed precision.



IDENTITY: MARK, CONSTRUCTION AND APPLICATIONS

FIG.07

a type drawing; it carries no orientation

