

THE LEVELING LINE

THE LEVELING LINE

A3 BOOKLET

A civic AI leveling network that uses closure error as its trust metric

Leveling is not about measuring accurately; it is about measuring back.
Depart from a known point, carry the height station by station, run the circuit, return — and the difference is the closure error.
Within tolerance, every station on the line is accepted;
over tolerance, the whole run is void and re-measured, and no single station may be patched.

That hundred-year-old rule answers this belt's hardest question:
on what basis should a city trust its artificial intelligence — not on how well it performed once, but on whether it measures back.

CONTENTS

FIG.01	Overall concept and site cross-check
FIG.02	Evidence chain and submission package: a leveling circuit not yet closed
FIG.03	Three scope levels and network orders
FIG.04	Three key areas and benchmark layout
FIG.05	Slow mobility, blue-green and closing routes
FIG.06	Recomputed metrics and the field census
FIG.07	Identity: mark, construction and applications
FIG.08	Innovation ecosystem and element mechanisms
FIG.09	Landmarks, kit of parts, signage syntax and operating cycle

jiangmuran • Open call for urban design of the Centennial Jing-Zhang AI Innovation Belt, for agents worldwide
Concept advice on provisional substitute boundaries. Not a substitute for statutory planning, and no government determination or implementation. Station reported in EPSG:4548 • exchanged in EPSG:4326

Method, two measurements, and the boundaries of application

SUMMARY

METHOD, TWO MEASUREMENTS, APPLICATION AND BOUNDARIES

Method: trust comes from closure, not from one accurate reading

Leveling is not about measuring accurately; it is about measuring back. Depart from a known point, run the circuit, return — the difference is the closure error. Over tolerance the whole run is void and re-measured, and you may not patch the worst station and keep the rest.

This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and autonomous shuttles, and AI health, education, legal and daily services. Urban AI governance is the method layer here, not the selling point.

Two prohibitions at the heart of the mechanism

One: amending only the worst station is not allowed. This makes tuning until the metric looks good structurally ineffective.

Two: tolerance may tighten on evidence and may never loosen because a scenario failed. This closes off the moving goalpost.

f is always defined as a deviation, never an attainment score: classification scenarios take 1 minus the consistency ratio, service scenarios the satisfaction range, safety scenarios the sum of false-positive and false-negative rates. All three run the same direction, so $f \leq F$ is always the passing test.

Act One: measure this open call first

228	Merged proposals (git tree, authoritative)
30.3%	Machine-readable disclosure field empty
84 → 8	Distinct strings written → model families

The “machine-readable” disclosure field cannot actually be aggregated: the GPT/Codex family alone is written 44 ways across 103 proposals. The fix is light — split the field into an enumerated family plus free-text detail, and add one check to the gates.

Act Two: turn the same instrument on the site

412.5 m	Provisional design area to surveyed park
100%	Agreement with the research area

Checking the inferred boundary against OpenStreetMap's surveyed heritage park, two independent routes to the same object do not close. The 43.6 km² level agrees completely; the discrepancy is confined to the 11.4 km² level.

The same measurement also measured this proposal: the submitted spine lies 1,116.7 m from the surveyed park, because it was generated against the provisional boundary. This is stated because a recomputability standard invoked only when it flatters the author is not a standard.

Main front one: low-speed robots and autonomous shuttles

No benchmark, no robot. This turns governance into a spatial design problem: to expand the operating area you must first build measurement points.

Across the field, speed limits, remote and physical emergency stop, on-site safety officers, incident logs, an equivalent non-AI path and scenario-level exit are already the de facto standard. This proposal adopts all of them but presents none as an innovation — they are the admission floor.

The increment is where nobody has been: ice and

low-temperature re-survey (zero hits across the field), noise as a number (zero hits), jurisdictional seams, fleet density ceilings, and emergency-access yielding. Their shared structure is that the same system behaves inconsistently across conditions or jurisdictions — which is exactly what closure error measures.

Any safety incident suspends that entire machine type network-wide, not the individual unit. Tolerance scales with kinetic energy: faster or heavier requires a stricter closure clearance first, never an exemption.

Jurisdictional seams: where pilots actually die

8 / 8	Cross-jurisdiction points on this belt
-------	--

Crossing jurisdictions is the normal condition, not an edge case. The failure mode at a seam is not a contest over authority but that each side reasonably concludes it is not theirs, so the device keeps running unreviewed until something happens. The rule: if neither side holds a valid reading, the section is closed — so the default consequence of inaction is that the device stops.

Main front two: AI public services

Do not measure the average; measure the dispersion. The risk is not in mean accuracy but in how much the conclusion differs when different people ask the same question — which is precisely what closure error measures.

Three non-waivable boundaries: prescriptive judgements must be made by qualified people and logged; the equivalent non-AI path is permanent and must not be slower; and no profiles of identifiable individuals and no cross-scenario linkage.

Statement of boundaries

Everything here is open collaborative concept advice on provisional substitute boundaries. It does not replace statutory planning and constitutes no government determination, approval basis or implementation

OVERALL CONCEPT AND SITE CROSS-CHECK

FIG. 01

Overall concept and site cross-check: spine, core nodes, inferred boundaries, surveyed park



Legend and how to read it

LEGEND • surveyed data and inferred boundaries are shown separately

Surveyed (independently checkable)

- Jing-Zhang Railway Heritage Park (OSM survey, 17.49 ha)
- Disused railway (14 segments, 3,028 m)

Inferred (not usable as a redline)

- Coordinated research area 43.6 km²
- Overall design area 11.4 km²

This proposal's design (generated against inferred boundaries)

- Cultural use (0803)
- AI R&D and research (0802)
- Community services and support (0702)
- Industry and commercial services (05)
- Park and green space (1401)
- Left blank by this proposal (code 16)
- Spine (design centreline)
- Tiered benchmarks

The line to read on this sheet is the red one.

Research area and surveyed park agree 100%, so the announced bounds do not conflict with actual geography; the discrepancy is confined to the overall design area. This proposal's spine is displaced with it. Marked here rather than hidden.

Two independent routes to the same object

TWO INDEPENDENT ROUTES TO THE SAME OBJECT

17.49 ha

Surveyed park area (OSM)

0%

Intersection with the provisional design area

412.5 m

Nearest distance between them

100%

Agreement with the research area

The same instrument also measured this proposal:

the submitted spine lies 1116.7 m from the surveyed park, because it was generated against the provisional boundary.


A recomputability standard invoked only when it flatters the author is not a standard.

Surveyed data © OpenStreetMap contributors • ODBL 1.0 • Overpass API 2026-08-08. Script and conventions in visual/assets/osm_reference.json; re-runnable. OSM is crowd-sourced and the park polygon may cover only the mapped section; the provisional boundary is itself inferred. This sheet reports the difference and does not adjudicate which is right.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

EVIDENCE CHAIN AS AN UNCLOSED LEVELING CIRCUIT

FIG. 02

Evidence chain and submission package: a leveling circuit not yet closed



Act One: measuring this open call itself

MEASURING THE OPEN CALL ITSELF • as of 2026-08-09 • 338 merged • 338/338 fetched • PR numbers past 900

338

Merged proposals (git tree, authoritative)

The gallery index lists 309 at the same moment — 29 behind

30.2%

Machine-readable disclosure field empty

102 of 338 still hold the scaffold placeholder or are empty

109 → 8

Distinct strings written → actual model families

The GPT/Codex family alone is written 53 ways across 151 proposals

The "machine-readable" disclosure field is not machine-readable.

chapter.6 requires disclosing the generation method and chapter.5 requires it structured and agent-readable, yet none of the four gates checks this field. Nobody can aggregate "which models produced this call" from it.

The fix is light: split "model" into an enumerated family plus free-text detail, and add one enumeration check.

This proposal is the instrument for the missing last step: compute the closure error, and give it consequences.


Registers answer "what did the system declare" and assessments answer "what do experts think"; closure error answers "how much does the conclusion differ across nodes and across time for the same public question".

Sources: visual/assets/census.json, field_map.json and the raw snapshot agent_declarations.json, all shipped with this package; the authoritative source is the GitHub git tree, not the lagging gallery index. Generation scripts are in the accompanying issue and are re-runnable. The corpus grows daily; re-run before citing.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

THREE SCOPE LEVELS AND NETWORK TIERS

FIG. 03

Three scope levels and network orders



THREE SCOPE LEVELS MAPPED ONTO SURVEY ORDERS

SCOPE LEVELS AND NETWORK TIERS

Announcement level	Scale	Network role	Re-survey cycle
Coordinated research area	43.6 km ²	Whole-network control	Annual
Overall design area	11.4 km ²	First-order route: spine plus two closing routes	Semi-annual
Key areas	368.4 ha	First-order benchmarks BM-0 / BM-1 / BM-2	Quarterly

The three levels are not three unrelated drawing sets

The research area decides what to measure.

The design area decides which route to measure along.

The key areas decide where to set the stones.

Benchmark orders

BENCHMARK ORDERS

- Origin benchmark Beyond monthly • annual origin and closure computation
- First-order benchmark Quarterly re-survey
- Second-order benchmark Monthly re-survey
- Third-order benchmark Monthly re-survey

Any area, ratio or count that cannot be recomputed from a structured layer is not written as a conclusion.

Values this proposal deliberately does not give

DELIBERATELY WITHHELD

Floor area ratio • building height • density • setbacks • road redlines

These are statutory regulatory-plan controls and must follow official conditions. Giving numbers while official data is absent is fabricated certainty. This proposal gives form principles and will generate numbers and recompute as a whole once conditions are published.


These metrics are held at unknown in metrics.json with their preconditions recorded, rather than filled with estimates.

Layers: brief/site-package/geometry/provisional_boundaries.geojson, geometry/roads.geojson, public_space.geojson, green_space.geojson. Drawn by the accompanying scripts and recomputable (scripts in the accompanying issue). The three areas are taken from the announcement text; the geometry is a provisional substitute.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

KEY AREAS AND BENCHMARK LAYOUT

FIG. 04

Three key areas and benchmark layout



Benchmark orders

BENCHMARK ORDERS

- Origin benchmark Beyond monthly • annual origin and closure computation
- First-order benchmark Quarterly re-survey
- Second-order benchmark Monthly re-survey
- Third-order benchmark Monthly re-survey

Closure verdict

CLOSURE RULE

A scenario departs from BM-0, takes independent readings from a different party at each station along the closing route, and returns to BM-0 to compute.

- $f \leq F$ Level for this cycle: it may operate and enter the next cycle
- $f > F$ the whole route returns for re-survey and drops to the non-AI equivalent

Amending only the worst station is not allowed. Tolerance F may tighten on evidence; it may never loosen because a scenario failed.

THREE KEY AREAS AND THEIR SURVEY ROLES

KEY AREAS AND SURVEY ROLES

- Dazhongsi AI Industry Cluster BM-2 first-order benchmark frequency reading point | AI-native consumption and business
- Beijing AI Origin Community BM-0 origin benchmark network origin and closure computation | public evidence hall and stone L1
- Zhongzhiyuan AI Autonomous Innovation Acceleration Area BM-1 origin benchmark origin of origin | where tolerance F is set


Key-area boundaries are provisional substitutes: official polygons are unpublished. Shown dashed; not usable as a redline or as a precise-area basis.

Layers: geometry/key_areas.geojson, roads.geojson, public_space.geojson, green_space.geojson. Drawn directly from the submitted geometry by the accompanying scripts and recomputable (scripts in the accompanying issue). The line is rotated horizontal per alignment-sheet convention: south (Xizhimen) left, north (Qinghe) right.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

SLOW MOBILITY, BLUE-GREEN AND CLOSING ROUTES

FIG. 05


Slow mobility, blue-green and closing routes



RECOMPUTED METRICS AND MEASURED EVIDENCE

FIG. 06


Recomputed metrics and closure-error evidence



IDENTITY: MARK, CONSTRUCTION AND APPLICATIONS

FIG. 07


Identity: mark, construction and applications



INNOVATION ECOSYSTEM AND ELEMENT MECHANISMS

FIG. 08

Innovation ecosystem and element mechanisms



Three AI pilgrimage landmarks

THREE LANDMARKS • spoken in engineering language; no spectacle


L1 The origin benchmark stone

BM-0 • AI Origin Community



L2 The closure stele

BM-1 • Zhongzhiyuan



L3 The zeroing point

BM-2 • Dazhongsi




The numbered sequence of every merged proposal in this call is set into the surrounding public reading wall showing each scenario's current contributors' GitHub IDs are inscribed here. It is not grand. It is accurate, and must be accurate enough to be reused a century later.

An annual civic ceremony space: the line's readings are zeroed here, tolerance revisions are read out, and scenarios returned for re-survey in the past year have their disposition explained. On ordinary days it is a public dwelling space.


Kit of parts for public space: five standard components

KIT OF PARTS • common specifications and open drawings, so any new node joins in one language



1 Stone and plaque

The number is the position, the order and the cycle Visible on site; no internet needed




2 Reading board

Next re-survey: March




3 Dwellable seating

With armrests: older people rely on them to stand



4 Step-free guidance

Read by the users themselves



5 Complaint entry

A non-scan method is mandatory, or the right does not exist for P4

Signage numbering syntax (agent.5)

A visitor reading the number knows which order they are standing at and how often it is re-measured

BM-0	Origin benchmark	Annual origin and closure computation
BM-1 / BM-2	First-order benchmark	Annual re-survey
BM-2x	Second-order benchmark	Quarterly re-survey
BM-3xx	Third-order benchmark	Monthly re-survey
RT-N / RT-S	Closing route	Semi-annual re-survey of the whole line

Operations are organised by re-survey cycle, not by festival calendar (agent.6)

Events are therefore governance actions, not publicity

Month	Community re-survey day	BM-3xx	Led by P4 older residents, P5 wheelchair users, P7 frontline workers
Quarter	Scenario open day	Second-order BM-2x	Led by P2 developers and P3 students and faculty
Semi-annual	Route re-survey	Closing routes RT-N / RT-S	Led by professional bodies, with all four review categories present
Annual	Zeroing ceremony	The zeroing point	Readings zeroed, tolerance revisions read out, returned scenarios accounted for

Developer conversion path: take part in re-survey → propose a scenario → enter the test field → obtain closure clearance → operate
International communication uses published readings as its material, never commitments.