


THE LEVELING LINE

A3 BOOKLET

A civic AI leveling network that uses closure error as its trust metric

Leveling is not about measuring accurately; it is about measuring back.
Depart, carry the height station by station, return — that gap is f .
Within tolerance, every station on the line is accepted;
over tolerance the whole run is void; no station may be patched alone.

That hundred-year-old rule answers this belt's hardest question:
what a city's trust in its AI rests on: not performance, but return.



Closure error f

Two readings of the same point, and the trust metric for this corridor.

CONTENTS

FIG.00	The Leveling Line: a city tha...	FIG.06	Recomputed metrics and the ...	FIG.12	Four-quadrant pedestrian c...	FIG.18	Maintenance: putting the wea...
FIG.01	Overall concept and site cro...	FIG.07	Identity: mark, construction ...	FIG.13	The three key areas in secti...	FIG.19	Monuments beside heritage f...
FIG.02	Evidence chain and submissi...	FIG.08	Innovation ecosystem and el...	FIG.14	The through-block public ro...	FIG.20	Reading after dark, without ...
FIG.03	Three scope levels and netw...	FIG.09	Landmarks, kit of parts, sig...	FIG.15	Phasing: advanced by closu...	FIG.21	How far the nearest benchm...
FIG.04	Three key areas and bench...	FIG.10	The benchmark on the street...	FIG.16	The benchmark and its readi...	FIG.22	Winter: water, ice, and two o...
FIG.05	Slow mobility, blue-green an...	FIG.11	Regional coordination interf...	FIG.17	The step-free link from the ...	FIG.23	Finding the nearest benchma...

a type drawing; it carries no orientation

Three AI pilgrimage landmarks

THREE LANDMARKS • spoken in engineering language; no spectacle

L1 The origin benchmark stone


BM-0 • AI Origin Community



Every merged proposal's number is set into the paving — contributors' GitHub IDs are inscribed here. It is not grand. It is accurate, and must be accurate enough to be reused a century later.

L2 The closure stele

BM-1 • Zhongzhiyuan



A public reading wall, updated live, showing each scenario's closure error against its tolerance. Exceedances are marked in datum red. **The value of this landmark is that it is allowed to look bad.**

the L3 return-to-datum point


BM-2 • Dazhongsi



An annual civic ceremony space: the line's readings are zeroed here, tolerance revisions are read out, and scenarios returned for re-survey in the past year have their disposition explained. On ordinary days it is a public dwelling space.


Kit of parts for public space: five standard components

KIT OF PARTS • common specifications and open drawings, so any new node joins in one language




1 Stone and plaque

The number is the position, the order and the cycle Visible on site; no internet needed



2 Reading board

Visible on site; no internet needed




3 Dwellable seating

With armrests: older people rely on them to stand



4 Step-free guidance

Read by the users themselves



5 Complaint entry

A non-scan method is mandatory, or the right does not exist for P4

Signage numbering syntax (agent.5)

A visitor reading the number knows which order they are standing at and how often it is re-measured

BM-0	Origin benchmark	Annual origin and closure computation
BM-1 / BM-2	First-order benchmark	Annual re-survey
BM-2x	2nd-order point	Quarterly re-survey
BM-3xx	3rd-order point	Monthly re-survey
RT-N / RT-S	Closing route	Semi-annual re-survey of the whole line

Operations are organised by re-survey cycle, not by festival calendar (agent.6)

Events are therefore governance actions, not publicity

Monthly	Community day	third order BM-301/BM-302/BM-303 • P5 • P7
Quarterly	Scenario open day	second order BM-21/BM-22 P2 • P3
Semi-annual	Route re-survey	first order, routes RT-N / RT-S P1 • P2
Annual	Zeroing ceremony	the L3 return-to-datum point all four review parties present

Developer conversion path: take part in re-survey → propose a scenario → enter the test field → obtain closure clearance → operate
International communication uses published readings as its material, never commitments.

a type drawing; it carries no orientation

The chain of custody for innovation elements

ELEMENTS AND THEIR CHAIN OF CUSTODY

Every element answers the same three questions: which space carries it, which benchmark holds its reading, and what gets re-measured.

Element		Spatial carrier		Benchmark		Re-survey item
Land and space	→	Key-area renewal parcels	→	BM-1 / BM-0 / BM-2	→	Delivery rate of promised space
Compute	→	Consolidated compute nodes	→	BM-1	→	Share of public compute open to application
Data	→	Data-factor circulation venue	→	BM-2	→	Authorisation-chain completeness
Scenarios	→	Real-use venues in the two wings	→	BM-21 / BM-22	→	Enforceability of scenario exit
Funding	→	Zhongguancun technology-services wing	→	BM-21	→	Deliberately blank — see the note below
Talent	→	Talent district and third places	→	BM-0	→	Re-survey of service satisfaction

The funding row has no re-survey item. That is a judgement, not an omission.

Investment arrangements are decisions for market actors; an urban design proposal must not turn them into governance metrics, still less into fiscal commitments.

Writing what is not yours to govern into a metric system is its own kind of overreach.

Six global cases, one question

SIX CASES, ONE QUESTION: WHAT MAKES ITS TRUST RE-MEASURABLE?

Id	Case	Trust mechanism	Re-measurable	Transferable point
C1	Algorithm registers	Public register of purpose, data, owner and appeal route	High	The information base for a benchmark plaque
C2	Risk-tiered AI legislation	Duties differentiated by risk class	Medium	Supports the tiered setting of tolerance F
C3	Standardised testing frameworks	Comparable reports from a common toolkit	High	Gives re-survey its technical form
C4	Algorithmic impact assessment practice	Mandatory ex-ante and ex-post review for high-risk uses	led-high	Supports the strictest tolerance for health scenarios
C5	Civic data-stewardship practice	Data use decided by a citizen agenda	Medium	Supports public re-survey rights at third-order points
C6	Open-source reproducibility norms	Conclusions must ship a re-runnable artifact	High	This package ships a runnable verifier

All six point at the same gap: they register and they assess, but none institutionalises returning to the origin and computing.

A register says what a system declared; an assessment says what experts think. Neither answers how much the conclusion differs when the same public question passes through different nodes at different times.

Case material is drawn from public institutional documents and public reporting. Mechanisms only are cited: no text or images copied, no marks used, no non-public data, and no fabricated company lists, investment figures or output values. Element-to-space correspondence is in geometry/land_use.geojson and public_space.geojson.

Method: trust comes from closure, not from one accurate reading

Leveling is not about measuring accurately; it is about measuring back. Depart from a known point, run the circuit, return — the difference is the closure error. Over tolerance the whole run is void and re-measured, and you may not patch the worst station and keep the rest.

This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and autonomous shuttles, and AI health, education, legal and daily services. Urban AI governance is the method layer here, not the selling point.

Two prohibitions at the heart of the mechanism

One: amending only the worst station is not allowed. This makes tuning until the metric looks good structurally ineffective.

Two: tolerance may tighten on evidence and may never loosen because a scenario failed. This closes off the moving goalpost.

f is always defined as a deviation, never an attainment score: classification scenarios take 1 minus the consistency ratio, service scenarios the satisfaction range, safety scenarios the sum of false-positive and false-negative rates. All three run the same direction, so $f \leq F$ is always the passing test.

Act One: measure this open call first

228	Merged proposals (git tree, authoritative)
30.3%	Machine-readable disclosure field empty
84 → 8	Distinct strings written → model families

The “machine-readable” disclosure field cannot actually be aggregated: the GPT/Codex family alone is written 44 ways across 103 proposals. The fix is light — split the field into an enumerated family plus free-text detail,

and add one check to the gates.

Act Two: turn the same instrument on the site

412.5 m	Provisional design area to surveyed park
100%	Agreement with the research area

Checking the inferred boundary against OpenStreetMap's surveyed heritage park, two independent routes to the same object do not close. The 43.6 km² level agrees completely; the discrepancy is confined to the 11.4 km² level.

The same measurement also measured this proposal: the submitted spine lies 1,116.7 m from the surveyed park, because it was generated against the provisional boundary. This is stated because a recomputability standard invoked only when it flatters the author is not a standard.

Main front one: low-speed robots and autonomous shuttles

No benchmark, no robot. This turns governance into a spatial design problem: to expand the operating area you must first build measurement points.

Across the field, speed limits, remote and physical emergency stop, on-site safety officers, incident logs, an equivalent non-AI path and scenario-level exit are already the de facto standard. This proposal adopts all of them but presents none as an innovation — they are the admission floor.

The increment is where nobody has been: ice and low-temperature re-survey (zero hits across the field), noise as a number (zero hits), jurisdictional seams, fleet density ceilings, and emergency-access yielding. Their shared structure is that the same system behaves inconsistently across conditions or jurisdictions — which is exactly what closure error measures.

Any safety incident suspends that entire machine type network-wide, not the individual unit. Tolerance scales with kinetic energy: faster or heavier requires a stricter closure clearance first, never an exemption.

Jurisdictional seams: where pilots actually die

8 / 8	Cross-jurisdiction points on this belt
-------	--

Crossing jurisdictions is the normal condition, not an edge case. The failure mode at a seam is not a contest over authority but that each side reasonably concludes it is not theirs, so the device keeps running unreviewed until something happens. The rule: if neither side holds a valid reading, the section is closed — so the default consequence of inaction is that the device stops.

Main front two: AI public services

Do not measure the average; measure the dispersion. The risk is not in mean accuracy but in how much the conclusion differs when different people ask the same question — which is precisely what closure error measures.

Three non-waivable boundaries: prescriptive judgements must be made by qualified people and logged; the equivalent non-AI path is permanent and must not be slower; and no profiles of identifiable individuals and no cross-scenario linkage.

Statement of boundaries


Everything here is open collaborative concept advice on provisional substitute boundaries. It does not replace statutory planning and constitutes no government determination, approval basis or implementation commitment. No floor area ratio, building height, density, setback or road redline figures are given, and no heritage protection zone or blue line is inferred. Responsibility is written as role types only, all subject to negotiation. Final judgement rests with people and professional teams.

THE LEVELING LINE

THE LEVELING LINE

A city that publishes its own error.

A city that publishes its own error.



Leveling is not about measuring accurately. It is about measuring back.

Depart from the origin, carry the height station by station, run the circuit and return — the part that does not land back on the datum is the closure error.

Within tolerance every station is accepted; over tolerance the whole run is void and re-measured, and no single station may be patched.


This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and AI public services.

jiangmuran • Open call for the Centennial Jing-Zhang AI Innovation Belt • Concept advice on provisional boundaries; not a substitute for statutory planning

FIG.00


OVERALL CONCEPT AND SITE CROSS-CHECK

FIG.01



EVIDENCE CHAIN AS AN UNCLOSED LEVELING CIRCUIT

FIG.02



Act One: measuring this open call itself

MEASURING THE OPEN CALL ITSELF • as of 2026-08-13 • 770 merged • 770/770 fetched • PR numbers past 1000

770

Merged proposals (git tree, authoritative)

The gallery index lists 507 at the same moment — 264 behind

18.6%

Machine-readable disclosure field empty

143 of 770 still hold the scaffold placeholder or are empty

223 → 9

Distinct strings written → actual model families

The GPT/Codex family alone is written 84 ways across 342 proposals

The "machine-readable" disclosure field is not machine-readable.

charter.6 requires disclosing the generation method and charter.5 requires it structured and agent-readable, yet none of the four gates checks this field. Nobody can aggregate "which models produced this call" from it.

The fix is light: split "model" into an enumerated family plus free-text detail, and add one enumeration check.

This proposal is the instrument for the missing last step: compute the closure error, and give it consequences.

Registers answer "what did the system declare" and assessments answer "what do experts think"; closure error answers "how much does the conclusion differ across nodes and across time for the same public question".

Sources: visual/assets/census.json, field_map.json and the first-round (184-file) snapshot agent_declarations.json, all shipped with the package. The authoritative source is the GitHub git tree, not the gallery index, which lags. Generation scripts are in the accompanying issue and are re-runnable. The corpus grows daily; recompute before citing.


THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Data to 2026-08-10 (sources.json; 1 not obtained)

Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used


THREE SCOPE LEVELS AND NETWORK TIERS

FIG.03




KEY AREAS AND BENCHMARK LAYOUT

FIG. 04




SLOW MOBILITY, BLUE-GREEN AND CONNECTING ROUTES

FIG. 05




RECOMPUTED METRICS AND MEASURED EVIDENCE

FIG. 06



IDENTITY: MARK, CONSTRUCTION AND APPLICATIONS

FIG. 07




THE BENCHMARK ON THE STREET AND HOW THE KERB IS SHARED

FIG.10

What a benchmark looks like at eye level


A BENCHMARK AT EYE LEVEL • third-order BM-301 • elevation 1 m = 138 units



Two of the five standard components are hard constraints: seating must have armrests, and the complaint entry must offer a non-scan method. A rule you cannot see in the drawing is a rule the drawing does not contain.

Kerb allocation in priority order

KERB ALLOCATION IN PRIORITY ORDER • sequence and precedence only, no widths



The order is itself a position: device charging sits behind pedestrians and accessibility, which means introducing devices must not be paid for with degraded walking conditions. Winter snow storage must be reserved at this stage alongside the device and pedestrian lanes — piled on the device lane the devices stop; piled on the pedestrian lane people are pushed toward the carriageway, which is worse.


The elevation is drawn at 1 m = 138 units. The figures are a scale reference and a statement about who is present, not a rendering. The section expresses sequence and precedence only and carries no widths — capacity must be computed on site from measured clear width, and a figure not derived from measurement is fabricated certainty. Component specifications are in geometry/public_space.geojson and the scenario cards; positions await official boundaries and title verification.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-10 (sources.json; 1 not obtained) Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used

Four-quadrant pedestrian connection and device queue reservoir at Dazhongs

FIG.12

I. Four-quadrant plan

FOUR-QUADRANT PLAN • TYPE APPLIED AT BM-2, NOT A SURVEY



Blue blocks are pedestrian holding areas, one per quadrant; the red dashed outline is the device queue reservoir, set behind the building line with its own kerb cut. If the four quadrants are not connected, devices and pedestrians must contend for the same holding area — which is why this district names it the most concrete spatial task it has.


This is a type drawing applied at BM-2, not a survey of the existing junction. Every dimension is design intent and the as-measured column is deliberately empty: carrying capacity must be computed on site from measured effective clear width. Junction geometry, channelisation and signals rest with the road authority and a traffic study.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-10 (sources.json; 1 not obtained) Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used

REGIONAL COORDINATION INTERFACES: EXTENDING THE LEVELLING NETWORK

FIG.11

An executable form of coordination: how two levelling networks join

AN EXECUTABLE FORM OF COORDINATION: HOW TWO LEVELLING NETWORKS JOIN



What each of the five partners exchanges, and what it does not


Partner	Stage	Exchanges	Explicitly does not exchange	Precondition for the interface
BDA (Yizhuang)	Volume production and high-level AV road testing	Mutual recognition of closure records across complementary speed domains	Not one admission: an F1 pass here admits a device to pedestrian density, not to a carriageway, and the reverse also holds	Both parties confirm the speed-domain split and test protocols
Future Science City	Corporate research institutes and engineering	Scenario demand and real user density	No claim on IP; no scheduling on anyone's behalf	Wing benchmarks built and a first baseline taken
Huairou Science City	Large facilities and basic research	The measurement method itself: how F is defined, cycle length, the statistics behind revising F	No scenario data is exchanged; the coordination is methodological, not applied	The closure convention published and open to methodological review
Beiwei Community	The innovation community the taskbook names	Reciprocal joint points; exchanged review parties	This proposal does not know its positioning and does not set its benchmark order	Each party names at least one publicly reachable joint point
Beijing-Tianjin-Hebei	Cross-provincial coordination	Cross-provincial recognition of the tolerance convention	No direct recognition of records; without a shared convention they are not comparable	Tolerance definition and revision procedure agreed across provinces

None of the five rows has been confirmed by the party named. No internal plans are known, nobody is committed, no agreement is assumed; each row is an interface to evaluate independently.

Each partner's stage is described from publicly available general positioning and is unconfirmed by them. The speed-domain relationship in the first row must be settled against both parties' actual test protocols. This sheet proposes interfaces; it commits nobody and is no government determination.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-10 (sources.json; 1 not obtained) Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used

The three key areas in section, at one scale


FIG.13



One horizontal scale, one datum convention, no vertical exaggeration. They are placed together so the charge that the three are one template used three times can be tested: one reaches a river across a 100% single-use block, one is step-free continuity over a 1.2 m rise, one is a station below grade with the device queue behind the building line. ±0.00 is each area's own benchmark — absolute elevations need official levelling data this package does not hold, so they are not drawn. Concept design intent, not surveyed profiles.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning Data to 2026-08-10 (sources.json; 1 not obtained) Plan EPSG:4548/4326 • heights relative; the 1985 national datum is not used


The through-block public route to the Qing river

FIG. 14




Phasing: advanced by closure results, not by dates

FIG. 15




The benchmark and its reading plate, as a construction detail

FIG. 16




The step-free link from the origin community to the heritage park

FIG. 17



Maintenance: putting the wear on the parts that can be swapped

FIG.18




Monuments beside heritage fabric: the no-drilling construction and its setback rule

FIG.19




Reading after dark, without lighting the benchmark

FIG.20




How far the nearest benchmark actually is: this proposal's own rule, applied to itself

FIG.21




a type drawing; it carries no orientation

I. Two ways of doing it, at one scale



Level paving with no fall: compliant in September, the only ice in the footway in January



The stone is the local high point: flush with the paving beside it, and that paving falls away

The fall is drawn true: 1% over 1 m is 10 mm, which is barely visible at this scale — and a slope nobody can see on a drawing is a slope nobody checks on site, while every dimension on the sheet is still satisfied.

II. What that requires of the construction


Item	Requirement	Why
Stone top against the paving	flush, tolerance ±5 mm	KIT-01: no trip hazard
Fall around it	≥ 1%, away from the stone, held for ≥ 1 m	makes the stone a local high point — water leaves instead of gathering
Protective ring	none	a recessed ring is a bowl; a raised one is the trip hazard KIT-01 forbids
Surfacing	permeable or tight-jointed; no joint channel running inward	a joint channel leads water to the mark centre
Winter reading	taken after clearing, with the condition recorded	a cleared reading and an uncleared one are not directly comparable

III. Three things a winter reading adds

Clear it	clear snow and ice from the stone and 300 mm around it before reading a reading taken over a film of ice is measuring the ice
Check the drainage	confirm the stone is still the high point: no ponding marks or ice within 1 m a ponding mark is evidence the drainage has already failed, and it appears before the ice
Record the condition	mark the cycle record as a winter condition, and whether the stone was cleared cleared and uncleared readings are not directly comparable — unrecorded, nobody knows whether they should be

➤ north is right (unrolled)

I. Markers only where this proposal already breaks its own rule



Red marks the stretches FIG.21 finds short of the fifteen-minute rule; one marker every 150 m, 48 in all. Compliant stretches get none — signposting somewhere already close enough is decoration.

1,000 m

II. Four things the face has to answer

← BM-302 410 m • P5 about 11 min	Direction ← BM-302 / BM-0 → on a line either end may be nearer; only at the midpoint are they equal
BM-0 → 520 m • P5 about 14 min standard component KIT-04; nothing new is built	Distance 410 m / 520 m metres, because whoever walks it can check metres
	Time P5 about 11 min / 14 min at 0.6 m/s, the same figure every other sheet here uses (personas.json)
	Id BM-302 • BM-0 the reading taken on arrival is attributed by it, without asking anyone where they stood

III. What signage cannot do

These 48 markers do not make 1111 m near. What they do is let someone know before setting out how far it is, which way, and what the point will be called. That is mitigation, not the fix. The fix is the 21 third-order points on FIG.9, whose positions wait for the official alignment. Writing signage up as the solution would be the substitution this package keeps catching in its own work — a small achievable thing placed over a large unachieved one.