


THE LEVELING LINE

A3 BOOKLET

A civic AI leveling network that uses closure error as its trust metric

Leveling is not about measuring accurately; it is about measuring back.
Depart, carry the height station by station, return — that gap is f .
Within tolerance, every station on the line is accepted;
over tolerance the whole run is void; no station may be patched alone.

That hundred-year-old rule answers this belt's hardest question:
what a city's trust in its AI rests on: not performance, but return.



Closure error f

Two readings of the same point, and the trust metric for this corridor.

CONTENTS

FIG.00	The Leveling Line: a city tha...	FIG.06	Recomputed metrics and the ...	FIG.12	Four-quadrant pedestrian c...	FIG.18	Maintenance: putting the wea...
FIG.01	Overall concept and site cro...	FIG.07	Identity: mark, construction ...	FIG.13	The three key areas in secti...	FIG.19	Monuments beside heritage f...
FIG.02	Evidence chain and submissi...	FIG.08	Innovation ecosystem and el...	FIG.14	The through-block public ro...	FIG.20	Reading after dark, without ...
FIG.03	Three scope levels and netw...	FIG.09	Landmarks, kit of parts, sig...	FIG.15	Phasing: advanced by closu...	FIG.21	How far the nearest benchm...
FIG.04	Three key areas and bench...	FIG.10	The benchmark on the street...	FIG.16	The benchmark and its readi...		
FIG.05	Slow mobility, blue-green an...	FIG.11	Regional coordination interf...	FIG.17	The step-free link from the ...		

Three AI pilgrimage landmarks

THREE LANDMARKS • spoken in engineering language; no spectacle

L1 The origin benchmark stone


BM-0 • AI Origin Community



Every merged proposal's number is set into the paving — contributors' GitHub IDs are inscribed here. It is not grand. It is accurate, and must be accurate enough to be reused a century later.

L2 The closure stele

BM-1 • Zhongzhiyuan



A public reading wall, updated live, showing each scenario's closure error against its tolerance. Exceedances are marked in datum red. **The value of this landmark is that it is allowed to look bad.**

L3 The zeroing point

BM-2 • Dazhongsi



An annual civic ceremony space: the line's readings are zeroed here, tolerance revisions are read out, and scenarios returned for re-survey in the past year have their disposition explained. On ordinary days it is a public dwelling space.

Kit of parts for public space: five standard components

KIT OF PARTS • common specifications and open drawings, so any new node joins in one language

BM-301
Third-order benchmark plaque
Monthly re-survey • unique number

1 Stone and plaque

The number is the position, the order and the cycle


Closure error f

0.14 Level, within tolerance

Next re-survey: March

2 Reading board

Visible on site; no internet needed



3 Dwellable seating

With armrests: older people rely on them to stand



4 Step-free guidance

Read by the users themselves

Scan to appeal
or phone / on-site registration

A number on receipt • reply within 5 working days
No number means not accepted, and that counts in the readings

5 Complaint entry

A non-scan method is mandatory, or the right does not exist for P4

Signage numbering syntax (agent.5)

A visitor reading the number knows which order they are standing at and how often it is re-measured

BM-0	Origin benchmark	Annual origin and closure computation
BM-1 / BM-2	First-order benchmark	Annual re-survey
BM-2x	2nd-order point	Quarterly re-survey
BM-3xx	3rd-order point	Monthly re-survey
RT-N / RT-S	Closing route	Semi-annual re-survey of the whole line

Operations are organised by re-survey cycle, not by festival calendar (agent.6)

Events are therefore governance actions, not publicity

Monthly	Community day	Third-order BM-301/302/303	Led by P4 older residents, P5 wheelchair users, P7 frontline workers
Quarterly	Scenario open day	Second-order BM-21/22	Led by P2 developers and P3 students and faculty
Semi-annual	Route re-survey	Closing routes RT-N / RT-S	Led by professional bodies, with all four review categories present
Annual	Zeroing ceremony	3 The zeroing point	Readings zeroed, tolerance revisions read out, returned scenarios accounted for

Developer conversion path: take part in re-survey → propose a scenario → enter the test field → obtain closure clearance → operate
International communication uses published readings as its material, never commitments.

The chain of custody for innovation elements

ELEMENTS AND THEIR CHAIN OF CUSTODY

Every element answers the same three questions: which space carries it, which benchmark holds its reading, and what gets re-measured.

Element		Spatial carrier		Benchmark		Re-survey item
Land and space	→	Key-area renewal parcels	→	BM-1 / BM-0 / BM-2	→	Delivery rate of promised space
Compute	→	Consolidated compute nodes	→	BM-1	→	Share of public compute open to application
Data	→	Data-factor circulation venue	→	BM-2	→	Authorisation-chain completeness
Scenarios	→	Real-use venues in the two wings	→	BM-21 / BM-22	→	Enforceability of scenario exit
Funding	→	Zhongguancun technology-services wing	→	BM-21	→	Deliberately blank — see the note below
Talent	→	Talent district and third places	→	BM-0	→	Re-survey of service satisfaction

The funding row has no re-survey item. That is a judgement, not an omission.

Investment arrangements are decisions for market actors; an urban design proposal must not turn them into governance metrics, still less into fiscal commitments.

Writing what is not yours to govern into a metric system is its own kind of overreach.

Six global cases, one question

SIX CASES, ONE QUESTION: WHAT MAKES ITS TRUST RE-MEASURABLE?

Numbering	Case	Trust mechanism	Re-measurable	Transferable point
C1	Algorithm registers	Public register of purpose, data, owner and appeal route	High	The information base for a benchmark plaque
C2	Risk-tiered AI legislation	Duties differentiated by risk class	Medium	Supports the tiered setting of tolerance F
C3	Standardised testing frameworks	Comparable reports from a common toolkit	High	Gives re-survey its technical form
C4	Algorithmic impact assessment practice	Mandatory ex-ante and ex-post review for high-risk uses	led-high	Supports the strictest tolerance for health scenarios
C5	Civic data-stewardship practice	Data use decided by a citizen agenda	Medium	Supports public re-survey rights at third-order points
C6	Open-source reproducibility norms	Conclusions must ship a re-runnable artifact	High	This package ships a runnable verifier

All six point at the same gap: they register and they assess, but none institutionalises returning to the origin and computing.

A register says what a system declared; an assessment says what experts think. Neither answers how much the conclusion differs when the same public question passes through different nodes at different times.

Case material is drawn from public institutional documents and public reporting. Mechanisms only are cited: no text or images copied, no marks used, no non-public data, and no fabricated company lists, investment figures or output values. Element-to-space correspondence is in geometry/land_use.geojson and public_space.geojson.

Method, two measurements, and the boundaries of application

Method: trust comes from closure, not from one accurate reading

Leveling is not about measuring accurately; it is about measuring back. Depart from a known point, run the circuit, return — the difference is the closure error. Over tolerance the whole run is void and re-measured, and you may not patch the worst station and keep the rest.

This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and autonomous shuttles, and AI health, education, legal and daily services. Urban AI governance is the method layer here, not the selling point.

Two prohibitions at the heart of the mechanism

One: amending only the worst station is not allowed. This makes tuning until the metric looks good structurally ineffective.

Two: tolerance may tighten on evidence and may never loosen because a scenario failed. This closes off the moving goalpost.

f is always defined as a deviation, never an attainment score: classification scenarios take 1 minus the consistency ratio, service scenarios the satisfaction range, safety scenarios the sum of false-positive and false-negative rates. All three run the same direction, so $f \leq F$ is always the passing test.

Act One: measure this open call first

228	Merged proposals (git tree, authoritative)
30.3%	Machine-readable disclosure field empty
84 → 8	Distinct strings written → model families

The “machine-readable” disclosure field cannot actually be aggregated: the GPT/Codex family alone is written 44 ways across 103 proposals. The fix is light — split the field into an enumerated family plus free-text detail,

and add one check to the gates.

Act Two: turn the same instrument on the site

412.5 m	Provisional design area to surveyed park
100%	Agreement with the research area

Checking the inferred boundary against OpenStreetMap's surveyed heritage park, two independent routes to the same object do not close. The 43.6 km² level agrees completely; the discrepancy is confined to the 11.4 km² level.

The same measurement also measured this proposal: the submitted spine lies 1,116.7 m from the surveyed park, because it was generated against the provisional boundary. This is stated because a recomputability standard invoked only when it flatters the author is not a standard.

Main front one: low-speed robots and autonomous shuttles

No benchmark, no robot. This turns governance into a spatial design problem: to expand the operating area you must first build measurement points.

Across the field, speed limits, remote and physical emergency stop, on-site safety officers, incident logs, an equivalent non-AI path and scenario-level exit are already the de facto standard. This proposal adopts all of them but presents none as an innovation — they are the admission floor.

The increment is where nobody has been: ice and low-temperature re-survey (zero hits across the field), noise as a number (zero hits), jurisdictional seams, fleet density ceilings, and emergency-access yielding. Their shared structure is that the same system behaves inconsistently across conditions or jurisdictions — which is exactly what closure error measures.

Any safety incident suspends that entire machine type network-wide, not the individual unit. Tolerance scales with kinetic energy: faster or heavier requires a stricter closure clearance first, never an exemption.

Jurisdictional seams: where pilots actually die

8 / 8	Cross-jurisdiction points on this belt
-------	--

Crossing jurisdictions is the normal condition, not an edge case. The failure mode at a seam is not a contest over authority but that each side reasonably concludes it is not theirs, so the device keeps running unreviewed until something happens. The rule: if neither side holds a valid reading, the section is closed — so the default consequence of inaction is that the device stops.

Main front two: AI public services

Do not measure the average; measure the dispersion. The risk is not in mean accuracy but in how much the conclusion differs when different people ask the same question — which is precisely what closure error measures.

Three non-waivable boundaries: prescriptive judgements must be made by qualified people and logged; the equivalent non-AI path is permanent and must not be slower; and no profiles of identifiable individuals and no cross-scenario linkage.

Statement of boundaries


Everything here is open collaborative concept advice on provisional substitute boundaries. It does not replace statutory planning and constitutes no government determination, approval basis or implementation commitment. No floor area ratio, building height, density, setback or road redline figures are given, and no heritage protection zone or blue line is inferred. Responsibility is written as role types only, all subject to negotiation. Final judgement rests with people and professional teams.

THE LEVELING LINE

THE LEVELING LINE

A city that publishes its own error.

A city that publishes its own error.



Leveling is not about measuring accurately. It is about measuring back.

Depart from the origin, carry the height station by station, run the circuit and return — the part that does not land back on the datum is the closure error.

Within tolerance every station is accepted; over tolerance the whole run is void and re-measured, and no single station may be patched.


This proposal applies that hundred-year-old rule where a wrong reading injures someone: low-speed robots and AI public services.

jiangmuran • Open call for the Centennial Jing-Zhang AI Innovation Belt • Concept advice on provisional boundaries; not a substitute for statutory planning

FIG.00

OVERALL CONCEPT AND SITE CROSS-CHECK

FIG.01



Legend and how to read it

LEGEND • surveyed data and inferred boundaries are shown separately

Surveyed (independently checkable)

- Jing-Zhang Railway Heritage Park (OSM survey, 17.49 ha)
- Disused railway (14 segments, 3,028 m)

Inferred (not usable as a redline)

- Coordinated research area 43.6 km²
- Overall design area 11.4 km²

This proposal's design (generated against inferred boundaries)

Blank-block edges follow the existing arterials; they are not a parcel subdivision or an ownership boundary

- Cultural use (0803)
- AI R&D and research (0802)
- Community services and support (0702)
- Industry and commercial services (05)
- Park / green (1401)
- Urban road land (1207)
- Left blank by this proposal (code 16)
- Spine (design centreline)
- Tiered benchmarks

The line to read on this sheet is the red one.

Research area and surveyed park agree 100%, so the announced bounds do not conflict with actual geography; the discrepancy is confined to the overall design area. This proposal's spine is displaced with it. Marked here rather than hidden.

Two independent routes to the same object

TWO INDEPENDENT ROUTES TO THE SAME OBJECT

17.49 ha

Surveyed park area (OSM)

0%

Intersection with the provisional design area

412.5 m

Nearest distance between them

100%

Agreement with the research area

The same instrument also measured this proposal:

the submitted spine lies 1116.7 m from the surveyed park, because it was generated against the provisional boundary. A recomputability standard invoked only when it flatters the author is not a standard.


Surveyed data © OpenStreetMap contributors • ODbL 1.0 • Overpass API 2026-08-08. Script and conventions in visual/assets/osm_reference.json; re-runnable. OSM is crowd-sourced and the park polygon may cover only the mapped section; the provisional boundary is itself inferred. This sheet reports the difference and does not adjudicate which is right.

THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326

EVIDENCE CHAIN AS AN UNCLOSED LEVELING CIRCUIT

FIG.02



Act One: measuring this open call itself

MEASURING THE OPEN CALL ITSELF • as of 2026-08-12 • 702 merged • 702/702 fetched • PR numbers past 1000

702

Merged proposals (git tree, authoritative)

The gallery index lists 507 at the same moment — 195 behind

18.7%

Machine-readable disclosure field empty

131 of 702 still hold the scaffold placeholder or are empty

207 → 9

Distinct strings written → actual model families

The GPT/Codex family alone is written 78 ways across 314 proposals

The "machine-readable" disclosure field is not machine-readable.

charter.6 requires disclosing the generation method and charter.5 requires it structured and agent-readable, yet none of the four gates checks this field. Nobody can aggregate "which models produced this call" from it.

The fix is light: split "model" into an enumerated family plus free-text detail, and add one enumeration check.

This proposal is the instrument for the missing last step: compute the closure error, and give it consequences.

Registers answer "what did the system declare" and assessments answer "what do experts think"; closure error answers "how much does the conclusion differ across nodes and across time for the same public question".


Sources: visual/assets/census.json, field_map.json and the first-round (184-file) snapshot agent_declarations.json, all shipped with the package. The authoritative source is the GitHub git tree, not the gallery index, which lags. Generation scripts are in the accompanying issue and are re-runnable. The corpus grows daily; recompute before citing.

THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326

THREE SCOPE LEVELS AND NETWORK TIERS

FIG.03



THREE SCOPE LEVELS MAPPED ONTO SURVEY ORDERS

SCOPE LEVELS AND NETWORK TIERS

Announcement level	Scale	Network role	Re-survey
Coordinated research area	43.6 km ²	Whole-network control	Annual
Overall design area	11.4 km ²	First-order route: spine plus two closing routes	Semi-annual
Key areas	368.4 ha	Origin BM-0 with first-order BM-1 / BM-2	Quarterly

The three levels are not three unrelated drawing sets

The research area decides what to measure.

The design area decides which route to measure along.

The key areas decide where to set the stones.

Benchmark orders

BENCHMARK ORDERS

- Origin benchmark Annual origin and closure check
- First-order benchmark Annual re-survey
- 2nd-order point Quarterly re-survey
- 3rd-order point Monthly re-survey

Land-use partition composition

LAND-USE PARTITION

Left blank by this proposal (code 16)	22 features • 7 classes • no overlap, no gap (verified point by point)	55.4% • 11 blocks
AI R&D and research (0802)		16.9% • 2 blocks
Park / green (1401)		12.2% • 1 blocks
Community services and support (0702)		8.6% • 1 blocks
Industry and commercial services (05)		6.3% • 2 blocks
Cultural use (0803)		0.6% • 4 blocks
Urban road land (1207)		0.1% • 1 blocks

Reserved land at 55.4% is a judgement, not a gap: tenure, plan and retain/demolish status are unpublished; this proposal does not subdivide for them.

Values this proposal deliberately does not give

Floor area ratio • building height • density • setbacks • road redlines

Statutory regulatory-plan controls. Publishing a figure with the official data absent would manufacture certainty; metrics.json keeps them unknown.


Layers: brief/site-package/geometry/provisional_boundaries.geojson, geometry/roads.geojson, public_space.geojson. Drawn by the accompanying scripts and recomputable (scripts in the accompanying issue). The three areas are taken from the announcement text; the geometry is a provisional substitute.

THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326


KEY AREAS AND BENCHMARK LAYOUT

FIG.04




SLOW MOBILITY, BLUE-GREEN AND CONNECTING ROUTES

FIG.05




RECOMPUTED METRICS AND MEASURED EVIDENCE

FIG.06




IDENTITY: MARK, CONSTRUCTION AND APPLICATIONS

FIG.07



What a benchmark looks like at eye level


A BENCHMARK AT EYE LEVEL • third-order BM-301 • elevation 1 m = 138 units



Two of the five standard components are hard constraints: seating must have armrests, and the complaint entry must offer a non-scan method. A rule you cannot see in the drawing is a rule the drawing does not contain.

Kerb allocation in priority order

KERB ALLOCATION IN PRIORITY ORDER • sequence and precedence only, no widths



The order is itself a position: device charging sits behind pedestrians and accessibility, which means introducing devices must not be paid for with degraded walking conditions. Winter snow storage must be reserved at this stage alongside the device and pedestrian lanes — piled on the device lane the devices stop; piled on the pedestrian lane people are pushed toward the carriageway, which is worse.

The elevation is drawn at 1 m = 138 units. The figures are a scale reference and a statement about who is present, not a rendering. The section expresses sequence and precedence only and carries no widths — capacity must be computed on site from measured clear width, and a figure not derived from measurement is fabricated certainty. Component specifications are in geometry/public_space/geojson and the scenario cards; positions await official boundaries and title verification.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

An executable form of coordination: how two levelling networks join

AN EXECUTABLE FORM OF COORDINATION: HOW TWO LEVELLING NETWORKS JOIN



A joint observation produces a number, not an intention:

Each network reads the same public question set; the difference is the cross-network closure error. Within a jointly agreed tolerance, the two networks' records may be mutually recognised; over it, each re-surveys, and no declaration of partnership substitutes for a reading.

Nobody gives up authority: a joint observation asks only that conventions be comparable, not that organisations merge.

What each of the five partners exchanges, and what it does not

Partner	Stage	Exchanges	Explicitly does not exchange	Precondition for the interface
BDA (Yizhuang)	Volume production and high-level AV road testing	Mutual recognition of closure records across complementary speed domains	Not one admission: an F1 pass here admits a device to pedestrian density, not to a carriageway, and the reverse also holds	Both parties confirm the speed-domain split and test protocols
Future Science City	Corporate research institutes and engineering	Scenario demand and real user density	No claim on IP; no scheduling on anyone's behalf	Wing benchmarks built and a first baseline taken
Huairou Science City	Large facilities and basic research	The measurement method itself: how f is defined, cycle length, the statistics behind revising F	No scenario data is exchanged; the coordination is methodological, not applied	The closure convention published and open to methodological review
Beiwei Community	The innovation community the taskbook names	Reciprocal joint points; exchanged review parties	This proposal does not know its positioning and does not set its benchmark order	Each party names at least one publicly reachable joint point
Beijing-Tianjin-Hebei	Cross-provincial coordination	Cross-provincial recognition of the tolerance convention	No direct recognition of records; without a shared convention they are not comparable	Tolerance definition and revision procedure agreed across provinces


None of the five rows has been confirmed by the party named.
No internal plans are known, nobody is committed, no agreement is assumed; each row is an interface to evaluate independently.

Each partner's stage is described from publicly available general positioning and is unconfirmed by them. The speed-domain relationship in the first row must be settled against both parties' actual test protocols. This sheet proposes interfaces; it commits nobody and is no government determination.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

Four-quadrant pedestrian connection and device queue reservoir at Dazhongs

I. Four-quadrant plan

FOUR-QUADRANT PLAN • TYPE APPLIED AT BM-2, NOT A SURVEY



2. Dimensions: design intent / as-measured

DIMENSIONS: DESIGN INTENT VS AS-MEASURED


Item	Design intent	As measured
Pedestrian holding area, per quadrant width => 3.0 m x depth => 4.0 m	≥ 12 m²	
Crossing effective clear width after deducting poles, covers and the furnishing strip	≥ 4.0 m	
Kerb ramp gradient Barrier-Free Environment Construction Law and current codes	≤ 1:12	
device queue reservoir set behind the building line, with its own kerb cut	≥ 18 m²	
Reservoir to holding area, clear distance must not cross any desire line	≥ 2.0 m	
BM-2 stone to pedestrian edge the stone is flush with the ground and is not a trip hazard	≥ 0.6 m	

The right-hand column is deliberately empty.

Capacity is computed on site from measured clear width. A number filled in without measuring is manufactured certainty. Completed before any device is admitted.

3. Kerb allocation section (order and intended widths only)

KERB ALLOCATION SECTION



Blue blocks are pedestrian holding areas, one per quadrant; the red dashed outline is the device queue reservoir, set behind the building line with its own kerb cut.
If the four quadrants are not connected, devices and pedestrians must contend for the same holding area — which is why this district names it the most concrete spatial task it has.

This is a type drawing applied at BM-2, not a survey of the existing junction. Every dimension is design intent and the as-measured column is deliberately empty: carrying capacity must be computed on site from measured effective clear width. Junction geometry, channelisation and signals rest with the road authority and a traffic study.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

The three key areas in section, at one scale

One horizontal scale • 220 m across each • no vertical exaggeration • ±0.00 is that area's own benchmark, not a shared absolute datum

Zhongzhiyuan AI Acceleration Zone


±0.00 = BM-1

192.9 ha • BM-1

dominant use R&D 100%

2 land-use classes inside

External arrival. The area is 100% single-use, so public access to the Qing river runs through the block or not at all — what this section sets out is a public route across the R&D frontage and onto the levee path.



Beijing AI Origin Community


±0.00 = BM-0

104.3 ha • BM-0

dominant use community services 94%

2 land-use classes inside

A 1.2 m grade change, and step-free continuity. The campus edge ends in steps, and P5 needs a continuous step-free route. The reason this section exists is to put a 1:12 ramp with a landing where the steps are.



Dazhongs AI Industry Cluster


±0.00 = BM-2

72.0 ha • BM-2

dominant use industry and commerce 99%

2 land-use classes inside

A station below grade and a queue behind the building line. The only section here with anything below ground, and the only one with a carriageway. The device queue reservoir sits behind the building line with its own kerb cut, clear of the crossing desire lines — the four-quadrant plan is FIG.12.



One horizontal scale, one datum convention, no vertical exaggeration. They are placed together so the charge that the three are one template used three times can be tested: one reaches a river across a 100% single-use block, one is step-free continuity over a 1.2 m rise, one is a station below grade with the device queue behind the building line. ±0.00 is each area's own benchmark — absolute elevations need official levelling data this package does not hold, so they are not drawn. Concept design intent, not surveyed profiles.
THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning
Recomputed in EPSG:4548 • exchanged in EPSG:4326

Maintenance: putting the wear on the parts that can be swapped

FIG.18

I. Order of removal: from what is swapped to what is never moved



One person is assumed to carry up to 23 kg. The stone, 400 × 400 × 600 mm of plain concrete at 2,400 kg/m³, works out at about 230 kg — the one part nobody can carry is exactly the one part that should never be replaced.

II. What fails, how often it is replaced, and whether it touches the datum

Part	How it fails	Replaced	Mass	Tool	Touches datum
Face panel of the reading plate	stale values, scratching, fading	checked each cycle, replaced when damaged	3.0 kg	hex key	no
QR and appeal sticker	the link dies, the surface wears	annual	0.2 kg	none	no
post	impact, corrosion	about 20 years, or after an impact	12.0 kg	socket wrench	no
seat panel	wear, graffiti	about 10 years	18.0 kg	socket wrench	no
Tactile guidance strip	abrasion, and repaving	about 15 years, with the paving	6.0 kg	paving tools	no
stone	impact, excavation, settlement	never replaced	230 kg	plant required	yes

III. After it fails: found → held → replaced → lost

Found	any re-survey that finds the face stale or missing records it in that cycle the plate's own rule — a stale value counts as not posted — is the detector
Held	the point is marked returned in datum red, shown at every third-order point in the area bad news belongs where it happened (errata E50)
Replaced	a standard part is swapped without touching the stone or its mark centre one person, one visit; heaviest part 18 kg
Stone lost	the height is re-established from the next order up; the closure record carries the break; the id is retired and never reused reusing an id would quietly join two different points into one history

IV. What these intervals add up to

Summing the intervals above across 8 benchmarks: about 1.42 replacements per point per year, roughly 11 across the line; at 0.5 staff-hours each that is about 6 hours a year. That sits inside the annual operating band already published, so it needs no separate budget line — maintenance is not this mechanism's hidden cost, it is already in the table.

Face panel of the reading plate	1.6/yr	
QR and appeal sticker	8.0/yr	
post	0.4/yr	
seat panel	0.8/yr	
Tactile guidance strip	0.5/yr	

Intervals and hours per swap are assumptions, not measurements; the stone is never replaced and is excluded.

A monthly re-survey is a maintenance problem first, and this package had said nothing about what fails, how often, or who carries the replacement. One asymmetry organises the sheet: the stone is the only part that must never be replaced and the only one a person cannot carry (about 230 kg) — a benchmark that has moved is not the same benchmark. So the wear goes on the swappable parts. Intervals and masses are estimates, not measured.


THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326

Monuments beside heritage fabric: the no-drilling construction and its setback rule

FIG.19

I. Choosing the construction by distance from the fabric



d = horizontal distance to the outer edge of the fabric




The rule is complete; its input is missing

Which line d is measured from depends on the officially published heritage protection area and its construction-control zone. This package does not hold that layer: each of the 3 shipped constraint features records it as a data gap in its own note, and this proposal does not infer it. So the sheet gives the rule, the thresholds and the three constructions, and leaves that line to whoever publishes it — guessing a protection boundary to make the drawing look complete would mean inventing the one thing here that is not this proposal's to decide.

II. Two constructions, one scale

Buried: foundation below the frost line



Surface-set: bonded, no drilling, no excavation



III. What each construction costs

Construction	Applies when	Foundation	Orders it may carry	Re-survey	What removal leaves
Buried stone	d ≥ 5 m	below the standard frost depth D + 100 mm	datum / first / second / third	per order	plant to lift out; leaves an excavation
set-back independent foundation	2 ≤ d < 5 m	below frost line, set back ≥ 2 m, unanchored	second / third	per order	plant to lift out; leaves an excavation
Surface-set marker (no penetration)	d < 2 m	no foundation; bonded to existing paving, no drilling	third order only	monthly, and tied to the nearest buried point each lifting	it leaves only a bond mark, which cleans off

The call is for a belt along a hundred-year-old railway; this proposal's central act is sinking concrete. Three constructions, chosen by distance from the fabric, with the cost stated: without excavation there is no foundation below the frost line, so the surface-set marker may not carry first- or second-order duty. Where d is measured from depends on the official heritage layer, absent here and recorded as a gap in constraints.geopjon.


THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326

Reading after dark, without lighting the benchmark

FIG.20

I. Geometry: where the light comes from, and returns to



II. Three approaches, compared by how they fail

Approach	Needs power	How it fails	Light trespass
Retroreflective face, reader brings the light	no	the face gets dirty → visible in daylight → caught before it stops being readable	none
Low-level floodlight	yes, at every point	the lamp dies → the plate looks intact and cannot be read	yes, needs controlling
Internally lit face	yes, at every point	the lamp dies → the whole face goes, and it still looks like a plate	yes, and aimed at the reader

The first is chosen not because it is cheap but because its failure is visible. When a lamp dies the plate still looks intact and can no longer be read — exactly the thing this proposal exists to prevent, a measurement whose failure cannot be seen. A dirty retroreflective face is simply dirty, visible in daylight to anyone walking past, well before it stops being readable.

plate face, raked 15° retroreflective film, returns light to its source
a specular surface would bounce it straight into the reader — hence not specular

III. What is not lit

No fixed luminaire	the reader brings the light; no power supply, no luminaire to maintain
The stone is not lit	it sits flush with the ground; lighting it adds no legibility, only night-time glare
Nothing aimed upward	no supplementary light may be emitted above the horizontal
Nothing crossing the boundary	the added vertical illuminance at neighbouring dwelling windows must meet the local limit — that figure comes from the local standard, which this package does not hold

Three third-order points are read monthly by residents, and much of the year those days fall after dark — and the package had never said how the plate is read at night. The obvious answer is a lamp, and it is the wrong one: it needs power, needs maintaining, throws light into the windows opposite, and when it dies the plate looks intact and cannot be read — a measurement whose failure is invisible, which is the thing this proposal exists to prevent. So the face is retroreflective with no fixed luminaire and the reader brings the light; a dirty face is simply dirty, in daylight.


THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326

How far the nearest benchmark actually is: this proposal's own rule, applied to itself

FIG.21

I. Distance to the nearest benchmark, along the spine



II. Turning that sentence into a spacing rule that can be checked

Where the rule comes from	this proposal's own repeated claim: a right of review you must walk fifteen minutes to reach has not been given
Interior spacing	chainage between adjacent points ≤ 2 × 540 m = 1080 m
Both ends of the line	end of line to its nearest point ≤ 540 m
Which speed the rule uses	PS's 0.6 m/s, not a typical adult's 1.2 m/s — setting the rule by the faster figure sets it at a standard most people cannot use

III. How far short the current layout falls

Six of the nine segments between the eight existing points fail. Closing the gap needs **9 more third-order points** (or 3 if the rule is only held to an average walker). They are not free: third-order points are read monthly, and every one of them enters the annual operating table. **This sheet does not place them** — putting 9 coordinates on a provisional substitute spine is the manufactured certainty this register counts. The rule ships; the positions wait for the official alignment.

This proposal argues repeatedly that a right of review you must walk fifteen minutes to reach has not been given, and its entire participation argument rests on anyone being able to walk up to a benchmark and take a reading — and nobody had ever measured that walk. Measured: the worst place on the line is 1107 m from the nearest benchmark, 5 minutes for P30.8, twice the limit this proposal holds others to. The rule ships as a spacing requirement; the positions wait for the official alignment.

THE LEVELING LINE • jiangmuran • Concept advice on provisional boundaries; not a substitute for statutory planning

Recomputed in EPSG:4548 • exchanged in EPSG:4326