



JUST ENOUGH JING-ZHANG

JUST ENOUGH JING-ZHANG

Right-Sized AI Urban Belt

算得更近，城市更稳 / Compute Closer, Keep the City Steady

CONCEPTUAL VISUALIZATION / 概念可视化

Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.



Right-Sized AI Urban Belt

Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.

PRI-01 · Minimum Sufficient Model

Prove that a task needs AI, then choose the smallest model, rule system, or combination that can complete it.
Could the task be done equally well without AI?

PRI-02 · Near-Edge First

When a device or site-level system can complete the task, raw data is not sent away by default.
Can the data remain where it is produced?

PRI-03 · Offline Usability

Wayfinding, information, help, and basic public services retain a complete offline or non-AI equivalent.
After disconnection, can the public continue without technical assistance?

PRI-04 · Graceful Degradation


Every enhanced service steps down smoothly from connected mode to local mode and then to a human or conventional service.
Is the degradation path continuous, visible, and rehearsable?

PRI-05 · Visible Resources

Show model size, update time, data scope, accountable owner, and conceptual performance targets in public language.
Can the public see what the system consumes and who is responsible?


PRI-06 · Human Final Judgment

Models never make the final decision on health, education, law, safety, planning, or the public interest.
Who can overrule the model, and where can the public appeal?





XingHE-YX · 2026-08-12 · 04



Zhongzhiyuan

Small-model proving field: compare baselines before models.

Beijing AI Origin

Offline co-creation field: no account, bilingual, accessible.

Dazhongsi

Near-intelligence adoption field: ownership, offline equivalent, funded exit.


Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.

==|==

SPACE FIRST, AI ENHANCES, FAILURE DOES NOT END PUBLIC SERVICE

04 / 05

NEAR-EDGE WAYFINDING · GREEN SPINE · THREE-LEVEL DEGRADATION · STAFFED ACCESS



THREE LEVELS FOR EVERY SERVICE

01 CONNECTED ENHANCEMENT

Updated content; raw data do not travel by default

02 LOCAL MODE

Signed content pack, local inference, no account

03 HUMAN / CONVENTIONAL

Paper map, mechanical sign, telephone, and staff

SPATIAL BASE: RAMP / TACTILE / CONTRAST / REST / WATER

OPERATING BASE: OWNER / UPDATE / DRILL / APPEAL / EXIT

CONCEPTUAL VISUALIZATION · PROVISIONAL ROUGH BOUNDARY (NOT AN OFFICIAL REDLINE) · AREAS RECALCULATED IN EPSG:4548 · 2026-08-12

01 Connected enhancement

Content updates; raw data do not travel by default.

02 Just-enough local

Signed content, local inference, no account.

03 Human / conventional

Paper maps, mechanical signs, telephones, and staff.

Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.



6 SMALL REVERSIBLE PROTOTYPES

BLDG-01 · Just Enough Zero Station test hall

Concept storeys: 2 · statutory height pending

BLDG-02 · Small-model repair workshop

Concept storeys: 2 · statutory height pending

BLDG-03 · Offline Co-Creation Court archive room

Concept storeys: 1 · statutory height pending

BLDG-04 · Public learning and accessibility test room

Concept storeys: 2 · statutory height pending

BLDG-05 · Near-Intelligence Market public-service desk

Concept storeys: 2 · statutory height pending

BLDG-06 · Offline emergency and operations station

Concept storeys: 1 · statutory height pending

4 PILGRIMAGE LANDMARKS

LM-01 · Just Enough Zero Station

A civic-task station that publicly compares no AI, a small model, and a remote model. Model-scale trials, rollback repair, and the six adoption gates

LM-02 · Offline Co-Creation Court

A public living room for design, storytelling, and deliberation without accounts or persistent connectivity. Bilingual stories, accessible navigation, and traceable public comment

LM-03 · Near-Intelligence Market

A visible civic-adoption front desk combining lightweight public services, daily repair, and no-account events. Public-service Q&A, microclimate rest, and emergency degradation

LM-04 · Compute Milestones

Small public-information nodes along the heritage park that make compute location, resource targets, updates, and ownership legible. Resource displays, offline maps, human help, and non-AI equivalents

Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.



Twelve Scenarios: Offline Equivalent and Exit for Every One

SCN-01 · Offline Multilingual Guide

A one-time-download guide to Jing-Zhang history, places, and public services that remains useful offline.
Offline: Paper bilingual maps, place numbers, and offline audio packs remain available.

SCN-02 · TEST · Accessible Near-Edge Navigation Trial

Tests personalized accessible routes against real surfaces, slopes, closures, and rest points.
Offline: Tactile paving, high-contrast signs, and human-verified accessible routes form the base system.

SCN-03 · TEST · Lightweight Public-Service Q&A Trial

Answers low-risk questions from a bounded policy corpus, showing a source and human service point for every response.
Offline: A searchable static FAQ, telephone list, and in-person service directory remain available.

SCN-04 · Microclimate Stay Advice

Turns site-level heat, humidity, wind, shade, and seating status into understandable stay options.
Offline: Shade maps, water points, rest points, and heat-warning signs remain visible.

SCN-05 · Facility Repair Assistant

Matches common asset identifiers to manual steps on a site terminal; staff confirm every work order.
Offline: Assets retain QR labels, paper quick guides, and duty phone numbers.

SCN-06 · Jing-Zhang Cultural Story Small Model

Generates age-appropriate bilingual short stories from edited historical material and shows the original archival trail on site.
Offline: Paper story cards, a timeline, and archive numbers work independently.

SCN-07 · TEST · On-Device Model-Scale Experiment

Lets participants compare rules and differently sized models on one public task for quality, latency, and conceptual resource targets.
Offline: Static baselines, paper scoring, and a model-free rule system serve as controls.

SCN-08 · Offline Emergency Information Board

Stores evacuation routes, assembly points, and emergency cards locally; connected mode only synchronizes signed updates.
Offline: Mechanical flip panels, photoluminescent maps, broadcasts, and human marshals remain in place.

SCN-09 · TEST · Low-Speed Robot Yielding Trial

Tests low-speed carriers yielding to children, wheelchairs, and cane users on a closed, visible loop.
Offline: The loop has physical separation, human observers, mechanical stops, and robot-free time windows.

SCN-10 · No-Account Event Reservation

Allocates public-event places through one-time ticket numbers and an on-site standby mechanism.
Offline: On-site ticketing, telephone registration, and a staffed standby list remain open.

SCN-11 · Small-Model Update and Repair Station

Provides signed, reversible, documented on-site updates for public terminals.
Offline: Every device retains the last verified version and a complete non-AI mode.

SCN-12 · Public-Comment Summarizer

Clusters comments that participants agreed to publish into traceable summaries, retaining minority views and links to originals.
Offline: A full comment index, human topic tags, and a paper comment wall remain open.



01 Connected enhancement

Updated content; raw data do not travel by default

Need low-cost real tasks, reproducible baselines, and public evaluation grounds. Verifiable tasks and clear adoption thresholds

02 Just-enough local

Signed content, near-edge inference, no account

Users

Need understandable services that do not depend on accounts, smartphones, or stable connectivity. One-step help and a visible human channel

03 Human / conventional

Paper maps, mechanical signs, telephones, and staff

Need safe, bounded learning and activity settings without collecting minors' data as the price of participation. Shared family use and no retention by default

PER-04 · Wheelchair and Sensory-Accessibility Users

Need continuous accessible routes, multisensory cues, and verifiable pre-trip information. Accessible space first; AI only as enhancement

PER-05 · International Developers, Visitors, and New Talent

Need bilingual guidance, portable interfaces, and trials that do not lock them to one platform. Bilingual, portable, and vendor-neutral access

PER-06 · Park and Facility Operators

Need maintainable systems with accountable ownership, rehearsed degradation, and a viable exit. Low maintenance burden and explicit shutdown conditions



PH-01 · Public Tasks and Baselines

Turn the concept into bounded tasks that communities and operators can verify together.
Months 0-6 after launch
A scenario enters trial only after it has a no-AI baseline, a data boundary, and a human owner.

PH-02 · Controlled Validation and Community Trial

Test public value and risk under observable conditions that can be stopped.
Months 7-18
Trials expand only after safety, accessibility, offline, and maintainability checks publish reproducible records.

PH-03 · Conditional Adoption and Long-Term Care

Move validated services into long-term operation with budgets, appeals, and exits.
Months 19-36 and beyond
Only services that consistently outperform the no-AI baseline and fund maintenance, appeal, and exit move into routine operation.

EV-01 · Spring Just Enough Problem Week

Residents, operators, and researchers turn daily frictions into testable public tasks.
Spring

EV-02 · Summer On-Device Experiment Season

Compares minimum-sufficient options under real heat, footfall, and disconnection drills.
Summer


EV-03 · Autumn Small-Model Open Day

Opens evaluations, failures, updates, resource targets, and adoption decisions to the public.
Autumn

EV-04 · Winter Offline Drill

Makes public services prove that they still work through disconnection, device failure, and staff handover.
Winter

OWNER · UPDATE · DRILL · APPEAL · EXIT



Recomputable

Areas, concept green and public-space ratios, prototype quantities, centre-line length.

Explicitly pending

Statutory FAR, height, density, setbacks, road redline and area.

Human judgment

Twenty-three tasks, nine standards, and fifteen depth items point to exact evidence.

Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.

SOURCE REGISTER

SITE-PACKAGE · Public Site Package and Provisional Extents

OFFICIAL-ANNOUNCEMENT · Open Call and Official Brief

AGENT-TASKBOOK · AI Agent Taskbook

BOUNDARY-SOURCE · Provisional Overall-Extent Source

KEY-AREA-SOURCE · Three Provisional Key-Area Extents

SRC-LOCAL-01 · Haidian 1+X+1 AI Innovation Layout

SRC-LOCAL-02 · 2026 Haidian District Government Work Report

SRC-LOCAL-03 · Jing-Zhang Railway Heritage Park Listing

SRC-LOCAL-04 · Jing-Zhang Railway Heritage Park Phase II Public-Space Information

SRC-LOCAL-05 · Haidian Urban-Renewal Action Guide

SRC-TECH-01 · MLPerf Tiny Benchmark

SRC-TECH-02 · LiteRT Overview

SRC-TECH-03 · ONNX Runtime Mobile

SRC-TECH-04 · ExecuTorch

SRC-TECH-05 · Web Neural Network API

SRC-GOV-01 · Artificial Intelligence Risk Management Framework

Current geometry is the repository's provisional rough extent, not an official redline. Areas support concept comparison only.

COPYRIGHT + AI GENERATION

Prose, structured data, cartography, charts, and layouts were generated and edited by Codex (OpenAI GPT-5 family) under XingHE-YX authorization. No peer-submission text or imagery, corporate logo, portrait, paid font, or third-party photograph is reproduced. Official notices, government pages, standards, and technical documents support fact checking and method only. The submission is CC-BY-4.0; external sources retain their rights. This concept does not replace statutory planning, engineering, government approval, or final human judgment.

16 sources · 7 assumptions · 23 tasks · 9 standards · 15 depth items